

SF-enhanced (SFe) 4

Technical Specification

Version 4.0-rc3 (23 January 2025)



Copyright © 2025 SFe Team and contributors

Copyright notice

All parts of this specification may be reproduced without additional written permission by the SFe Team, provided that you follow the license in section 2.3.

The SFe Team is in no way sponsored by or otherwise affiliated with Creative Technology Ltd. or E-mu Systems, Inc.

SoundFont® is a registered trademark of Creative Technology Ltd. The SoundFont® 2.04 file format is copyright © 1994–2002 E-mu Systems Inc.

To avoid confusion with the official SoundFont standard, the file format is referred throughout to as "SF" and/or "SFe". Future versions of SFe may not be compatible with legacy SF players.

That being said, "soundfont" is rapidly becoming a genericised word via the actions of the VGM (video game music) and lightsaber communities. If you are a member of the VGM, lightsaber, or another community, by not appropriating "SoundFont" for your unrelated formats, you can help keep the meaning of the trademark.

Any excerpts from SFSPEC24.PDF are copyrighted by Creative Technology Ltd, and are only used for reference. However, we are confident that the base file format is not copyrightable.

Disclaimers

This specification is subject to change without notice. Please obtain the latest version from the SFe Team GitHub page at <https://github.com/sfe-team-was-taken>.

This specification assumes familiarity of the SoundFont 2.04 file format (SFSPEC24.PDF), which can be found at <https://freepats.zenvoid.org/sf2/sfspec24.pdf>.

Table of contents

Section 1 Introduction.....	11
1.1 Preamble.....	12
1.2 Changelog.....	12
1.3 History of improvements to the SF format.....	13
1.4 Scope and purpose.....	14
1.5 Important differences from the legacy SF specification document.....	14
1.6 Document organisation.....	15
Section 2 Credits and copyright.....	17
2.1 Credits.....	18
2.2 License.....	18
Section 3 Versioning and updates.....	21
3.1 Specification Versioning.....	22
3.2 Future improvements.....	25
3.3 Long term support of SFe 4.....	25
Section 4 Terminology.....	27
4.1 Data structure terminology.....	28
4.2 Synthesis terminology.....	30
4.3 Parameter terminology.....	31
Section 5 SFe file format structure.....	33
5.1 File format extensions.....	34
5.2 General RIFF-type format structures.....	34
5.3 Chunk header types.....	35
5.4 RIFF error checking features.....	36
5.5 Structure of the SFe 4 file format.....	36
5.5.1 SFe 4 file format structure outline.....	36
5.5.2 ISFe-list information.....	38
5.5.3 Changed and removed chunks.....	38
5.5.4 String encoding.....	39
5.6 INFO-list chunk.....	40
5.6.1 ifil sub-chunk.....	40
5.6.2 Versioning rules.....	40
5.6.3 Specification versions to ifil values.....	40
5.6.4 isng sub-chunk.....	41
5.6.5 List of sound engines.....	41
5.6.6 ICRD sub-chunk.....	42
5.6.7 INAM, IENG, IPRD, ICOP, ICMT and ISFT sub-chunks.....	42
5.6.8 irom and iver sub-chunks.....	42
5.6.9 SFty sub-chunk.....	43
5.6.10 SFvx sub-chunk.....	44
5.6.11 flag sub-chunk.....	46
5.7 sdta-list chunk.....	47
5.7.1 smpl sub-chunk.....	47
5.7.2 SFe Compression.....	48

5.7.3 sm24 and sm32 sub-chunk.....	50
5.7.4 Using 8-bit samples.....	52
5.7.5 Looping rules.....	53
5.8 pdta-list chunk.....	54
5.8.1 phdr sub-chunk.....	54
5.8.2 New Bank System.....	54
5.8.3 inst sub-chunk.....	58
5.8.4 shdr sub-chunk.....	58
5.8.5 Other sub-chunks.....	58
Section 6 SFe enumerations and feature flags.....	61
6.1 About SFe enumerations.....	62
6.2 Feature flags.....	62
Section 7 Parameters and synthesis model.....	73
7.1 About the synthesis model.....	74
7.2 MIDI functions.....	74
7.3 Parameter units, generators, modulators and NRPNs.....	74
7.4 On implementation accuracy.....	74
Section 8 SFe error handling.....	77
8.1 Structurally Unsound errors.....	78
8.2 Non-critical errors.....	78
8.3 Duplicated preset locations within files.....	79
8.4 Duplicated preset locations across files.....	80
8.5 Undefined chunks.....	80
8.6 Unknown Enumerators.....	80
8.7 Maximum File Size Limit Exceeded.....	81
8.8 MIDI Errors.....	81
8.9 Illegal parameter values, out of range values, missing required items and illegal enumerators.....	81
Section 9 SiliconSFe.....	83
9.1 SiliconSFe overview.....	84
9.2 Header format.....	84
9.2.1 About the header format.....	84
9.2.2 romRiffHeader.....	85
9.2.3 romByteSize.....	85
9.2.4 romInterleaveIndex.....	85
9.2.5 romRevision.....	85
9.2.6 romVer.....	85
9.2.7 bankChecksum.....	85
9.2.8 bankChecksum2sComplement.....	85
9.2.9 bankSFeVersion.....	86
9.2.10 bankProduct.....	86
9.2.11 sampleCompType.....	86
9.2.12 bankStyle.....	86
9.2.13 bankCopyright.....	86
9.2.14 romSFeBankStart.....	86
9.2.15 romSineWaveStart.....	86

9.2.16 sampleSineWave.....	87
9.3 AWE ROM emulator.....	87
9.3.1 Introducing the AWE ROM emulator.....	87
9.3.2 ROM emulator sample specification.....	88
Section 10 Compatibility information.....	99
10.1 Specification and structural compatibility.....	100
10.2 Sample compatibility.....	102
Section 11 Program specification.....	105
11.1 Program compatibility levels.....	106
11.1.2 Sample specifications.....	108
11.1.3 Instrument specifications.....	108
11.1.4 Preset specifications.....	109
11.1.5 Player specifications.....	110
11.2 Converting between legacy SF and SFe.....	112
11.3 Converting between chunk header types.....	113
11.4 File repair programs.....	114
11.5 Why these guidelines?.....	126
11.6 How to test your program with SFSpecTest.....	130
11.7 Courtesy actions.....	131
Section 12 Glossary.....	133

Section 1

Introduction

1.1 Preamble

The SFe standard has been created to provide a successor to E-mu Systems®'s SoundFont® 2.04 standard.

- Large files (above 4 GiB) must use a 64-bit chunk header format.
- Programs that are designed to create existing SF2 files can easily be adapted to 64-bit with SFe.
- 64-bit headers will allow creators of sound banks to replace split banks with a large number of SF2 files, with large monolithic banks, where all samples and instruments can be used in every preset.
- SFe compatible players require a certain standard of features, which is in the program specification in section 11. This improves compatibility of SFe banks with players.

1.2 Changelog

Revision	Date	Description
4.0-rc3	23 January 2025	n/a

For draft specification revision history, see `draft-revision-history.md` (available in the SFe specification package or on the GitHub repository).

1.3 History of improvements to the SF format

In 2005, E-mu released SoundFont® 2.04 (specification date September 2002), the last version before abandoning the format. It included 24-bit support, and was designed for the Sound Blaster® X-Fi. Creative Technology Ltd is still around, but their current sound cards (such as the AE-9) do not support the SoundFont® format.

At an unknown time, Kenneth Rundt, the author of SynthFont and a series of products based on it, added some custom features, mostly ported over from the DLS format:

- Always play the sample to the end
- Velocity to volume envelope attack (from DLS)
- Velocity to modulation envelope attack (from DLS)
- Vibrato lfo to volume (from DLS)

Werner Schweer found a way to compress SoundFont® 2 files (with the lossy Vorbis format) around 2010.

- The resulting format was known as SF3.
- It is commonly used by open source programs, such as MuseScore and FluidSynth.
- This is the reason why this format is not called SF3.
- SFe incorporates the Werner SF3 specification as SFe Compression.

Another development was done by Cognitone, which created an open source program that can losslessly compress SoundFont® 2 files (with FLAC). This was done in 2017.

- This was unofficially called SF4.
- This is intended to be the true version 4, as Cognitone SF4 seems to not have been as widespread as SF3.
- SF4 has also been rejected by FluidSynth as too loosely defined.

Finally, stgiga found out that many programs don't mind RIFF64 (RF64) files.

- RIFF64 files have a much larger size limit than 32-bit RIFF files.
- It gives us a new horizon to experiment with longer, higher quality samples.
- There are also many other advantages of the RIFF64 format.

The SFe project started in 2020 as a proposal for a successor to legacy SF2.04 by Polyphone developer and SFe team member davy7125. Starting off as "SF3", it was renamed to avoid confusion with Werner SF3, firstly to "SF32 and SF64" but eventually to "SFe".

1.4 Scope and purpose

This is the specification for the SFe 4.0 file format, based on the famous E-mu Systems® SoundFont® 2.04, and Werner SF3 standards, and is intended to be a source of information on SFe.

To create files that support SFe, you will need:

- The SFe technical specification (this document)
- E-mu's provided documents for legacy SF2.04 ([SFSPEC24.PDF](#))

All features from legacy SF2.04 will be retained.

1.5 Important differences from the legacy SF specification document

- This is not a standalone document; it refers to [SFSPEC24.PDF](#). All relevant information from the Werner SF3 specification is included.
- This document is unofficial and was not created by E-mu. If they want this to be removed, we will remove it, but we hope that this is because they have their own official update to the SF format ready.
- For copyright reasons, we cannot copy information from the original standard, except for what is required to interpret this document and what is allowed under fair use.
- The SFe specification includes many additional requirements over reading the file format.
- Document organisation has been significantly modified to accommodate the intricacies of the SFe format, and to make it easier to read.

1.6 Document organisation

The sections in this specification are different to the sections in SFSPEC24.PDF, however they roughly correspond to some of these sections:

- Sections 1-3 of this specification to sections 0-1 of SFSPEC24.PDF
- Section 4 of this specification to section 2 of SFSPEC24.PDF
- Section 5 of this specification to sections 3-7 of SFSPEC24.PDF
- Section 6 of this specification to section 8 of SFSPEC24.PDF
- Section 7 of this specification to section 9 of SFSPEC24.PDF
- Section 8 of this specification to section 10 of SFSPEC24.PDF
- Section 9 of this specification to section 11 of SFSPEC24.PDF
- Section 12 of this specification to section 12 of SFSPEC24.PDF

Section 6 also contains feature flags along with enumerators, section 9 also contains information about the AWE ROM emulator that can be implemented, and sections 10 and 11 include information on compatibility concerns and guidance on writing SFe-compatible software respectively.

Significant differences from the previous version of the specification will be highlighted.

Section 2

Credits and copyright

2.1 Credits

SFe Team

Organisation

- GitHub: [SFe Team · GitHub](#)

Members

- sylvia-leaf (they/them) | Discord – @tanukilvia | Email – sylvialeaf6284@gmail.com
- Stgiga (they/them) | Discord – @stgiga
- spessasus
- davy7125

Want to join the SFe Team? Please contact sylvia-leaf using the above contact.

Special thanks

Thanks to these people or groups:

- derselbst (for format suggestions)
- mawe42 (for format suggestions)
- sagamusix (for feature suggestions)
- Werner Schweer (for creating Werner SF3, of which SFe Compression is based on)
- Falcosoft (for feature suggestions that will arrive in a future version of SFe)
- E-mu Systems (for creating the legacy SoundFont format, of which SFe is based on)

2.2 License

Copyright © 2024 SFe Team and contributors

Permission is granted to use, distribute and modify this specification for any use, provided that:

- you attribute the SFe Team (do not remove copyright notices)
- you clearly mark any modifications that are made to the specification
- you do not claim that we're affiliated with E-mu or Creative Labs.

THIS SPECIFICATION IS PROVIDED "AS-IS" WITHOUT ANY WARRANTY WHATSOEVER, INCLUDING IMPLIED WARRANTY.

Section 3

Versioning and updates

3.1 Specification Versioning

Final specifications have version numbers in the format x.yL, where x and y are numbers and L is a letter:

- x is incremented when a change in the SFe format is made in a way that makes the resulting files incompatible with the previous version.
- y is incremented when there are new features added to the SFe format.
- SFe should not skip "y" versions.
- L is incremented when there are small changes made to the specification. The first version will not include L.
- Release candidates have very similar version numbers to final specifications, but include "rc" between y and L.
- An example of a final specification version would be 4.0.

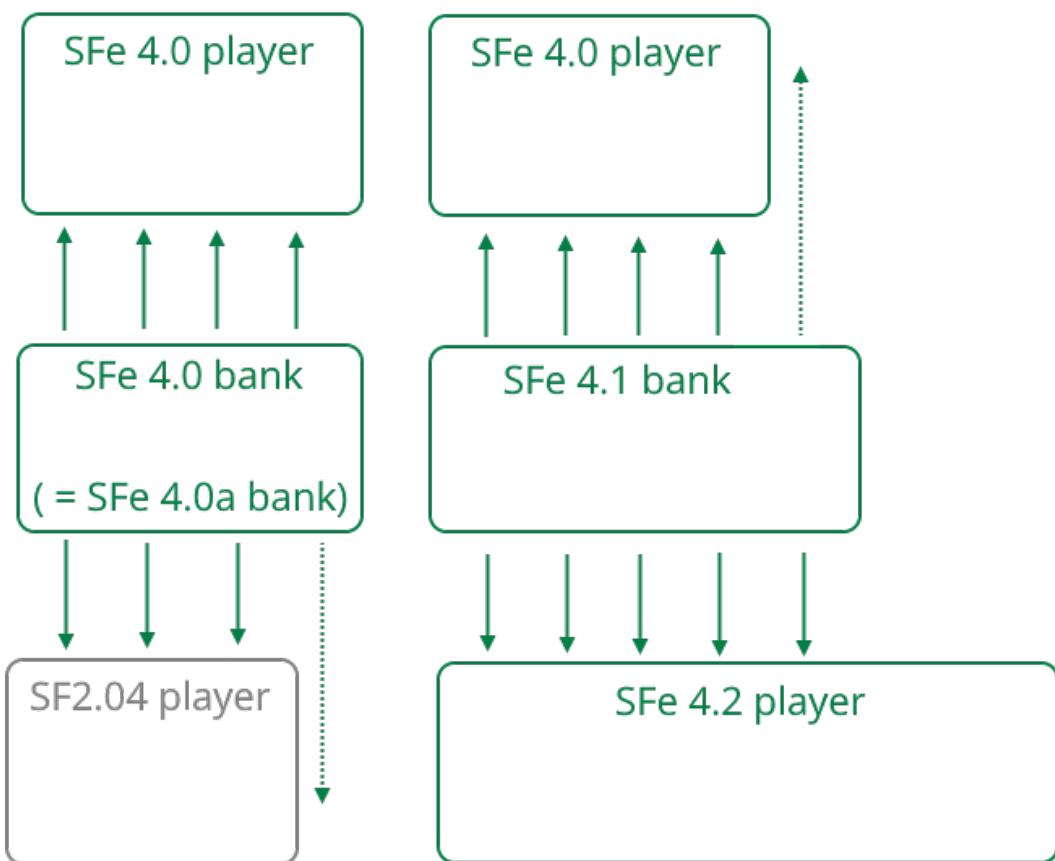


Figure 1: SFe 4.x versions are compatible with legacy players but with reduced sound quality. No changes are made to actual file structure in "L" versions. Later "y" versions are compatible with earlier players but with reduced sound quality.

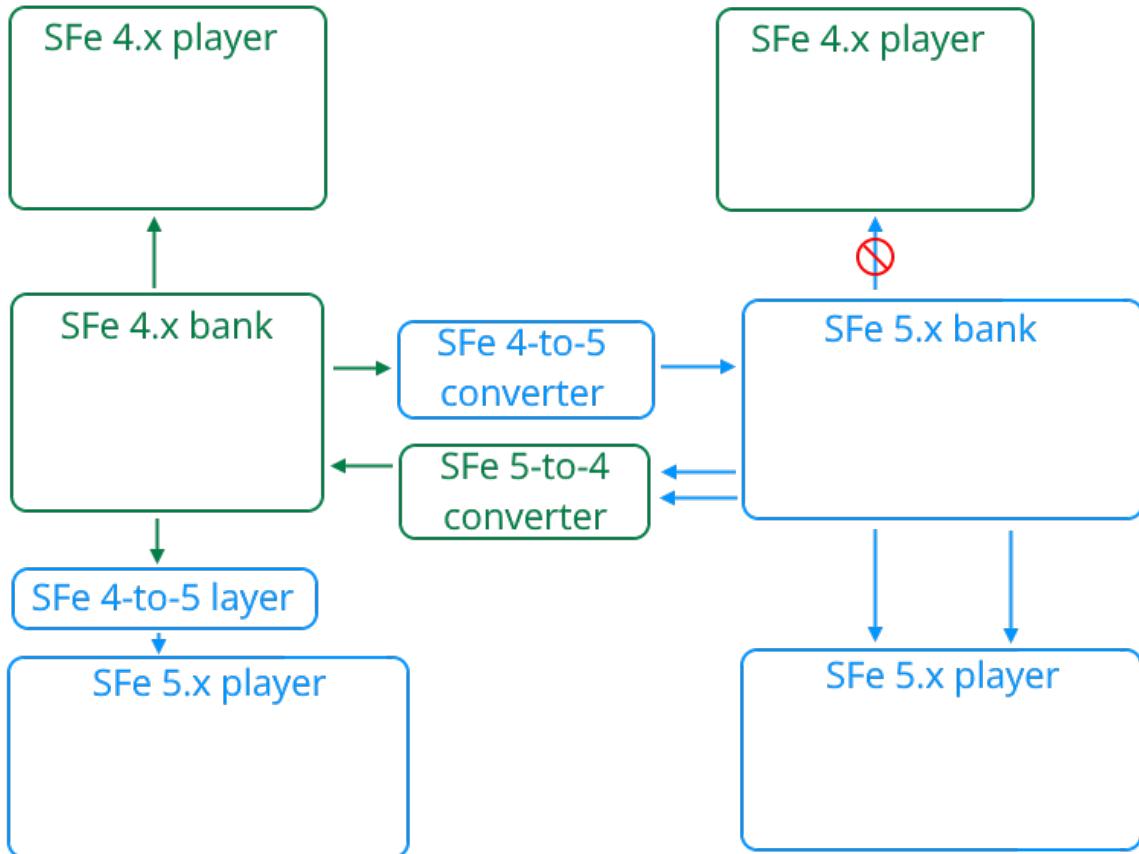


Figure 2: SFe "x" versions are not compatible with each other, but can be converted, resulting in reduced sound quality. Compatibility layers can be added to newer players to use older banks. It is also possible to upgrade banks to newer versions.

Draft specification milestones have version numbers in the format $x.y.zL$, where x , y and z are numbers and L is a letter. In this case, the versioning works similarly to a final specification, but with these changes:

- z is incremented when the draft undergoes a larger change, or large updates are made to the software.
- L is incremented when there are small changes, but only when pointed out by others. The first version will not include L .
- An example of a draft specification version would be 4.0.11a.

During the development of specifications, version numbers will be in the format $x.y.aaaabbccL$, where x , y , z , a , b and c are numbers, and L is a letter. The versioning is similar to final specifications and milestone drafts, but $aaaabbcc$ is the day in which the specification was updated, and L is incremented when updated.

Only final specifications are included in the changelog in section 1.2.

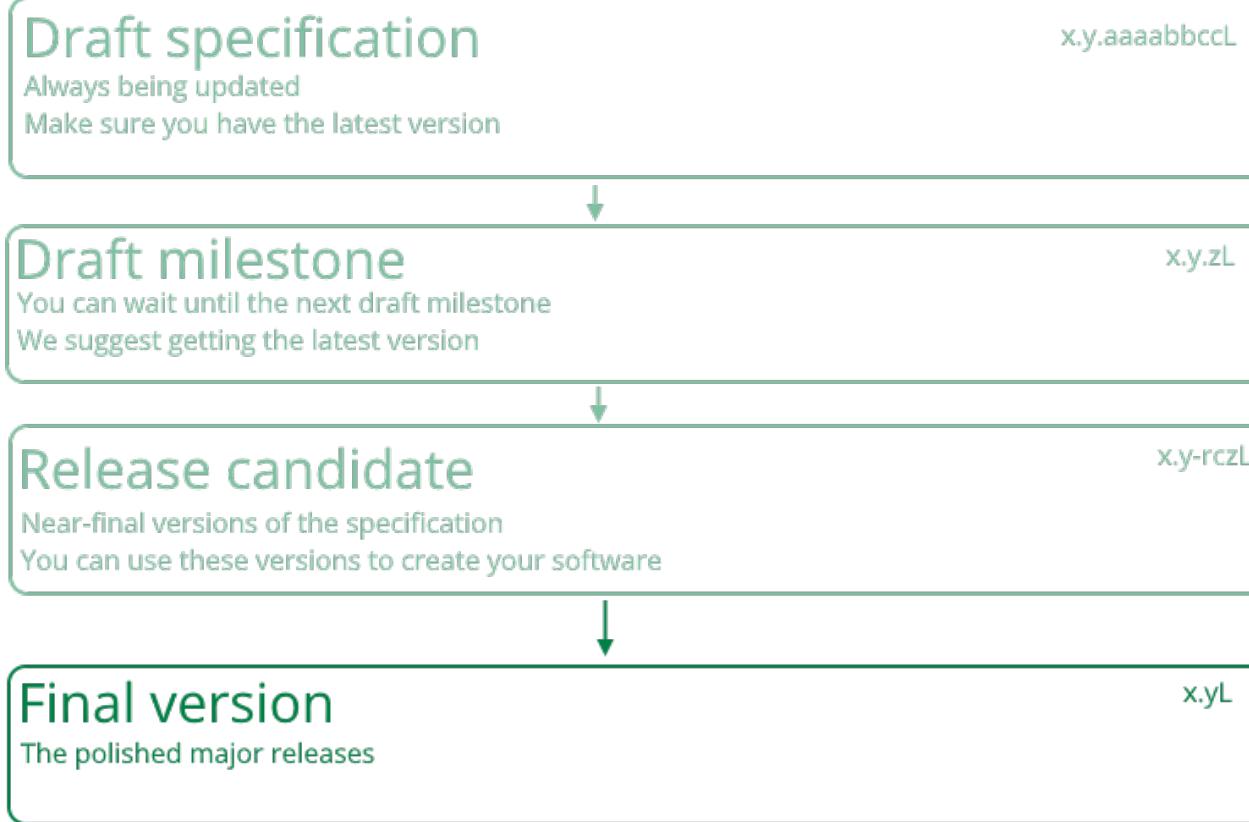


Figure 3: The SFe development process.

3.2 Future improvements

SFe is designed for future improvements.

- These will be done in a more liberal way than the conservative manner of the SoundFont® 2 updates that E-mu has done.
- SFe is not affected by limitations created by the EMU8000 sound processor and Sound Blaster® cards, and therefore will have more improvements.
- Additionally, starting from SFe 5, SFe won't be affected by limitations created by legacy SF players.
- To avoid over-stress of developers of the SFe Team, as well as SFe instrument banks, features will be spread out across versions, and similar features will arrive at the same time.

Here are a few things that are planned for SFe:

- Polyphone 3 will be the first program that supports SFe. It will use it by default with legacy SF being an option. (Polyphone 2.5 seems to be planned as an LTS release for legacy SF.)
- Negotiations with more SF program developers such as FluidSynth and Bassmidi will start soon.
- For SFe 4.1, there will be an overhaul of the default modulators system, inspired by the [DMOD proposal by spessasus](#). Support for the related PNMM sub-chunk will also be included.
- A MIDI lyrics specification for MIDI players, along with Spessasus/Falcosoft RMIDI support, will become available in SFe 4.2.
- We will negotiate with the Synthfont author Kenneth Rundt about getting the Synthfont Custom Features added for SFe 4.2. Care has been taken to ensure that SFe parameter usage does not conflict with SFCF.

3.3 Long term support of SFe 4

SFe 4 is a "long-term support" version, and will get feature updates along with later versions such as SFe 5. While this is the case, some features are structurally incompatible with legacy SF2.04, and will not be available in SFe 5.

If an earlier major version of SFe is the main version used for a longer time than expected, then it can be declared as another LTS version. Such a LTS version will be declared in this section.

Section 4

Terminology

4.1 Data structure terminology

The data structure terminology used in SFe 4.0 is broadly the same as legacy SF2.04, with these additions:

- Branch – A subdivision of a tree structure containing either sub-branches or leaves that include values.
- BW64 – Broadcast Wave 64, used in the RF64 Header.
- Cognitone SF4 – An incompatible modification to Werner SF3 to allow support for FLAC audio compression. Because it is considered proprietary compression, usage is not allowed in SFe.
- FLAC – A lossless audio compression format commonly used in open-source software. Supported by Werner SF3, but not commonly used for that purpose.
- Leaf – A value found in a tree structure at the end of a branch.
- Lossless compression – Said of a compression format that retains all of its data when compressed. In terms of audio, there is no loss in quality in losslessly compressed audio.
- Lossy compression – Said of a compression format that does not retain all of its data when compressed. In terms of audio, there is a loss in quality in lossily compressed audio.
- OGG – See "Vorbis".
- Opus – A lossy audio compression format, slightly newer than OGG but with less wide adoption.
- Proprietary Compression – Any non-Werner SF3 SoundFont® compression system. Usage is not allowed in SFe.
- RF64 – See "RIFF64".
- RIFF64 – A 64-bit RIFF-type format. Contrast to the RIFF format, which is 32-bit. Therefore, the maximum file size is above 4 gigabytes in size.
- RIFF-type format – Formats similar to RIFF (Resource Interchange File Format), see "RIFF" in SFSPEC24.PDF for more information.
- SFe – A family of enhancements to the SoundFont® 2.04 formats, unofficially created after E-mu/Creative abandoned the original format. May not be structurally compatible with legacy SF2.04.

- SFe 4 – This new specification, based on SoundFont® 2.04 and Werner SF3, with a set of new features making it more realistic. Not to be confused with the incompatible Cognitone SF4 file format.
- SFe-compatible – Indicates files, data, synthesisers, hardware or software that conform to the SFe specification.
- SFe Compression – The compression system based on Werner SF3 that SFe programs should be compliant with.
- Static RIFF – Any RIFF-type format with a fixed chunk size field width, including RIFF or RIFF64. See "RIFF-type format", "RIFF" and "RIFF64".
- Tree structure – A structure consisting of branches and leaves.
- Version 3 – See "Werner SF3".
- Version 4 – See "SFe 4".
- Vorbis – A lossy audio compression format commonly used in open-source software. The basic compression format that most Werner SF3 and SFe-compatible software should be expected to implement.
- Werner SF3 – A small upgrade to SoundFont® 2.04 created by Werner Schweer to allow an open source compression solution for SoundFont® programs. Standardised as SFe Compression.

And these changes:

- RIFF – The 32-bit static RIFF format used by SoundFont® 2.04 and the unmodified version of the format described by Microsoft in 1991.

4.2 Synthesis terminology

The synth terminology used in SFe 4.0 is broadly the same as legacy SF2.04, with these additions:

- AWE64 – The successor to the famous AWE32, adding features such as waveguide synthesis. Used the EMU8000 synthesizer chip, like the preceding AWE32. Available in "Value" or "Gold" versions.
- DAHDSR – Stands for Delay, attack, hold, decay, sustain, release. The six-step envelope system used in SF and SFe.
- EMU10K1 – The successor to the EMU8000, designed by E-mu® for the Creative Labs SB Live!.
- EMU10K2 – An update to the EMU10K1, designed by E-mu® for the Creative Labs SB Audigy.
- EMU20K1 – The successor to the EMU10K2, designed by E-mu® for the Creative Labs SB X-Fi.
- EMU20K2 – An update to the EMU20K1, please refer [here](#) for information on SB X-Fi cards that include it.
- Hold – The portion of the DAHDSR envelope after the attack portion, but before the decay portion starts.
- Legacy sound card – A Sound Blaster® (or other sound card) that uses a hardware MIDI synthesiser capable of using banks in the SoundFont® format.
- ROM samples – Obsolete feature used in legacy sound cards, most modern SF2 files do not use this feature.
- Sharp – Said of a tone that is higher in pitch than another reference tone.
- SB – Abbreviation of "Sound Blaster®". For example, "SB X-Fi".
- Sound Blaster® Live! – The successor to the AWE64, which improved the synthesizer chip to the EMU10K1, supporting modulators.
- Sound Blaster® Audigy – The successor to the SB Live!, containing the EMU10K2 chip.
- Sound Blaster® X-Fi – The successor to the SB Audigy, containing the EMU20K1 or EMU20K2 chip. Supports 24-Bit SoundFont® 2 files (2.04).
- Synth – Abbreviation of "Synthesiser," see "Synthesiser" in SFSPEC24.PDF for more information.

These changes:

- Articulation – Modulation of available parameters and usage of extra samples to produce expressive musical notes.
- Case-insensitive – Indicates that a UTF-8 character or string treats alphabetic characters of upper or lower case as identical.
- Case-sensitive – Indicates that a UTF-8 character or string treats alphabetic characters of upper or lower case as distinct.
- Downloadable – legacy SF2.0x, Werner SF3 or SFe file obtained from the internet. (Old meaning referred to the obsolete ROM system)
- MIDI Bank – Groups of up to 128 presets, which can be selected by the two MIDI "Bank Select" control changes (CC00 and CC32).

And these removals:

- Predator (refers to an old legacy SF2.0x editor made by E-mu)

4.3 Parameter terminology

The parameter terminology used in SFe 4.0 is broadly the same as legacy SF2.04, with these additions:

- Amplification – An increase in volume or amplitude of a signal.
- Flat – Said of a tone that is lower in pitch than another reference tone.

Section 5

SFe file format structure

5.1 File format extensions

The file format extension to use for SFe files is generally `.sft`:

- `.sf2` is avoided because SFe files are *not* SoundFonts, but simply files that use formatting that is very similar to legacy SF2.04.
- `.sf3` is avoided because some Werner SF3 bank players may not support SFe features.
- `.sf4` is avoided due to incompatibility with cognitone formatted banks.

When opening a bank with extension `.sft`, programs must determine the correct version to use. The `ifil` value and `ISFe-list` sub-chunk provide hints.

The file type should be referred to as SFe bank and should *not* be referred to as SoundFont or anything containing SoundFont. SFSPEC24.PDF states that files with additional chunks don't conform to the legacy SF2.04 standard.

SFe currently does not use a MIME type.

5.2 General RIFF-type format structures

RIFF-type formats are the file format used in legacy SF2.04, Werner SF3 and SFe standards.

There are a few different RIFF-type format structures:

- RIFF is the basic version with 32-bit chunk headers, and is used in legacy SF2.04 and Werner SF3.
- RIFF64 (also called RF64) is mostly compatible with RIFF, but uses 64-bit chunk headers.
- RIFX is a big-endian version of 32-bit RIFF, while RIFF and RIFF64 are little-endian formats.

RIFF-type formats are created in building blocks known as "chunks."

Chunks are defined using this structure:

- `ckID`: type of data in chunk, equal to a unique four character code (FourCC), listed above.
- `ckSize`: size of chunk (RIFF, RIFX), equal to 4,294,967,295 (RIFF64)
- `ds64`: size of chunk (RIFF64 only)
- `ckDATA[ckSize]`: the data inside the chunks, including pad bytes.

Chunks can be further divided into "sub-chunks."

Orders of chunks in all SF files are strictly defined, as in versions 2 and 3, and should be kept to, except for TSC mode.

5.3 Chunk header types

In SFe, there are different chunk header types that are used in the format. These correspond to different RIFF-type formats. Currently, there are two defined chunk header types:

- 32-bit static
 - This is the same as legacy SF
 - This corresponds to RIFF.
 - The FourCC used is RIFF.
- 64-bit static
 - A ds64 chunk is added
 - The FourCC used is RF64.
 - To prevent loading by incompatible players, the sfbk fourcc is replaced with sfen (**SF-enhanced**)
 - This corresponds to RIFF64.

Future versions of SFe may define different chunk header types.

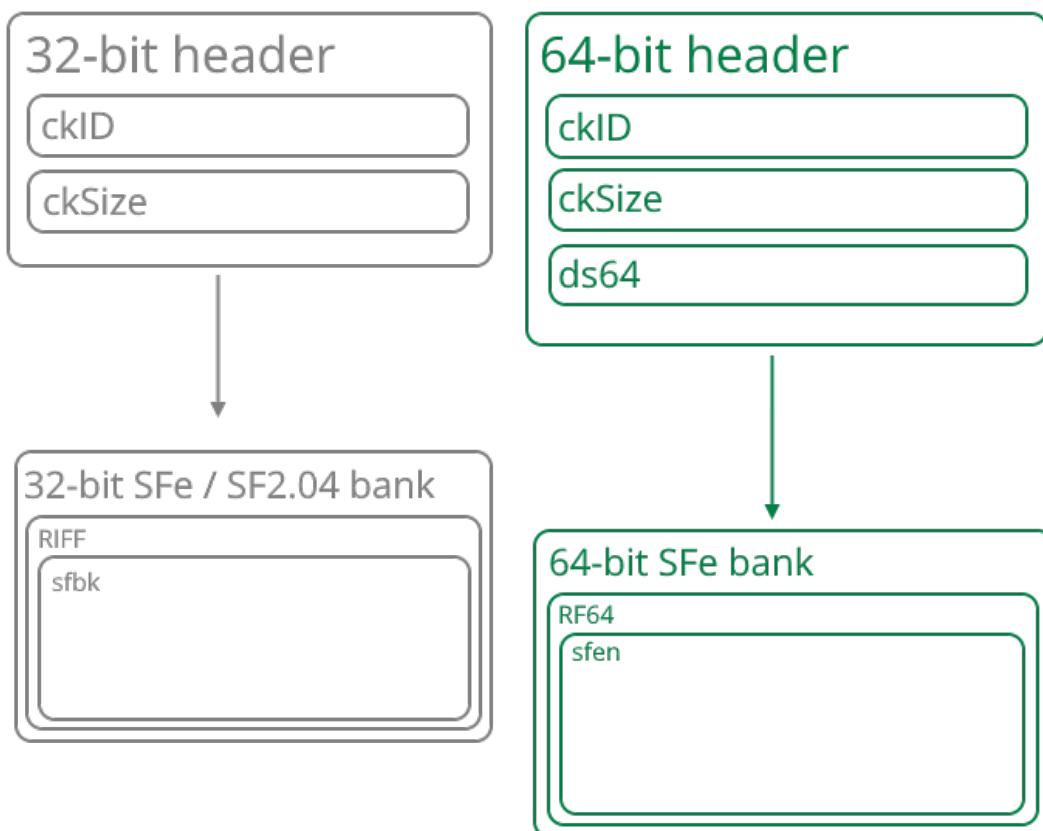


Figure 4: 32-bit static versus 64-bit static headers.

5.4 RIFF error checking features

RIFF-type formats have error checking features about:

- The size of the file
- The length of the chunks
- The length of the sub-chunks

Using this information, it is possible to check for damage to an SF(e) file.

5.5 Structure of the SFe 4 file format

5.5.1 SFe 4 file format structure outline

An SFe 4 file consists of:

- RIFF chunk (main chunk) - this changes depending on the chunk header type to be used.
 - sfbk ascii string – use sfen with 64-bit chunk headers
 - LIST
 - INFO ascii string
 - Sub-chunks inside INFO-list in legacy SF2.04 – ifil, isng, etc.
 - LIST
 - ISFe ascii string
 - Chunks listed in section 5.5.2
 - LIST
 - sdt_a ascii string
 - Sub-chunks inside sdt_a-list in legacy SF2.04 – smpl, sm24
 - sm32 chunk (BYTE array)
 - LIST
 - Sub-chunks inside pdta-list in legacy SF2.04

Only SFe-specific chunks are listed for brevity. In this section, assume that any non-listed chunk is identical to legacy SF2.04.

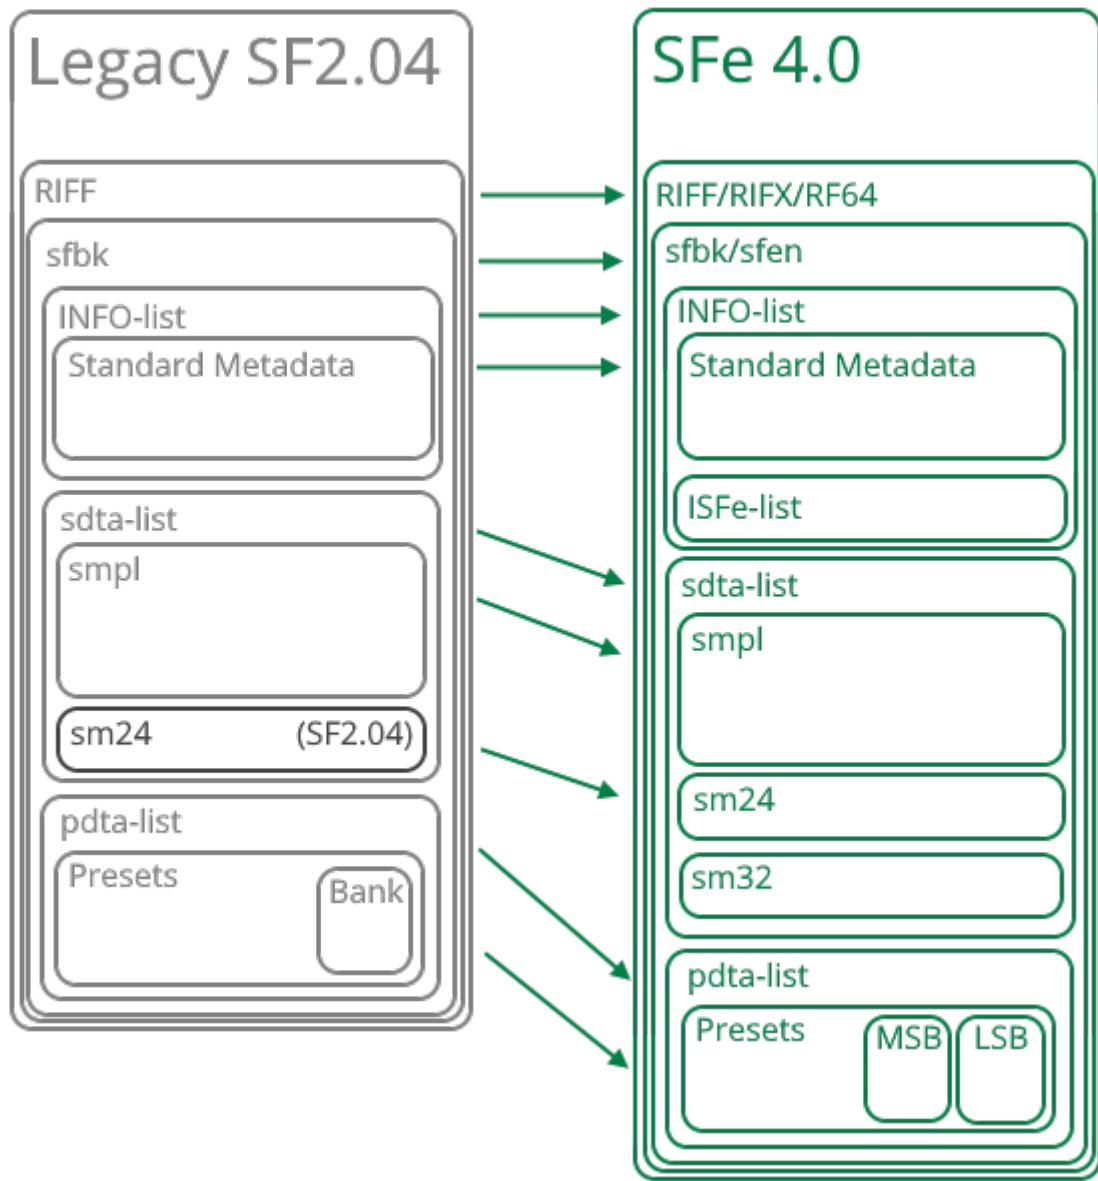


Figure 5: Legacy SF2.04 vs SFe 4.0 structures.

5.5.2 ISFe-list information

The ISFe-list sub-chunk includes many different sub-chunks to show information about SFe-specific features. Generally, we use the ISFe-list sub-chunk to make it clearer that this kind of information is SFe-specific.

Due to possible compatibility constraints, the ISFe-list sub-chunk is found inside the INFO-list sub-chunk, rather than as a fourth RIFF chunk. At least a few legacy soundcards (notably the SB X-Fi) do not error out on the inclusion of a fourth separate chunk. Officially, according to SFSPEC24.PDF, additional sub-chunks mean that an SFe file is not conformant to the legacy SF2.04 standard.

The ISFe-list sub-chunk currently contains these sub-chunks as of version 4.0:

- SFty chunk (UTF-8 string)
- SFvx chunk (46 bytes)
 - wSFeSpecMajorVersion (WORD)
 - wSFeSpecMinorVersion (WORD)
 - achSFeSpecType [20] (CHAR)
 - wSFeDraftMilestone (WORD)
 - achSFeFullVersion [20] (CHAR)
- flag chunk (multiple of 6 bytes)
 - byBranch (BYTE)
 - byLeaf (BYTE)
 - dwFlags (DWORD)

5.5.3 Changed and removed chunks

In the pdta-list sub-chunk, the wBank chunk has been replaced by byBankMSB and byBankLSB. These are functionally the same, but expressed in a different way to make the specification more readable.

5.5.4 String encoding

For most string fields, the encoding to use is now UTF-8 instead of ASCII. Mojibake may result on legacy SF players when using characters unsupported by ASCII. Because some characters use multiple bytes in UTF-8, you may not be able to use as many multi-byte characters compared to single-byte characters.

This applies to the `isng`, `INAM`, `irom`, `ICRD`, `IENG`, `IPRD`, `ICOP`, `ICMT`, `ISFT` chunks from legacy SF2.04, as well as the `achPresetName` (`PHDR`), `achInstrumentName` (`INST`) and `achSampleName` (`SHDR`) fields.

5.6 INFO-list chunk

5.6.1 ifil sub-chunk

The value of the `ofil` sub-chunk is equal to 2.1024 or 3.1024 depending on whether SFe Compression is used.

The size must be exactly four bytes. Reject files with an ifil sub-chunk that isn't four bytes as Structurally Unsound.

If the ifil sub-chunk is missing, either:

- Assume version 4.0.
- Reject the file as Structurally Unsound.

5.6.2 Versioning rules

In SFe 4.0, new versioning rules are used to replace the old ones.

The value of wMajor increases every time a change is made to the format that makes it incompatible with existing players.

- There will be at least 6 months between the first draft milestone of a new wMajor version and the release of the final specification.
- Older wMajor versions must be supported, either directly or via translation to the latest version.
- We strive to minimize the number of these updates whenever possible in favor of updates that are backward compatible.

The value of wMinor increases every time a change is made to the SFvx sub-chunk while retaining backwards compatibility.

- These updates generally have only one or two draft milestones before the final specification releases.
- Feature updates in these versions are smaller.

The specification type used is found in the ISFe-list sub-chunk.

5.6.3 Specification versions to ifil values

wSFeSpecMajorVersion	wSFeSpecMinorVersion	ofil (wMajor)	ofil (wMinor)
4	0	2 or 3	1024

5.6.4 isng sub-chunk

A new default isng sub-chunk value is used in SFe: SFe 4.

- SFe 4.0 players should recognize this and remove the default velocity related filter used in legacy SF2.04.
- In the case of a missing isng chunk, files with an ifil sub-chunk with wMajor = 2 or 3 and wMinor >= 1024, or wMajor >= 4, assume an isng sub-chunk value of SFe 4. Don't assume EMU8000.

Additionally, UTF-8 is now used instead of ASCII, and the length limit is removed.

- The isng sub-chunk contains a UTF-8 string of any length.
- Example of value: SFe 4 (with appropriate zero bytes).

Reject anything not terminated with a zero byte, and assume the value SFe 4. Do NOT assume EMU8000 by default.

5.6.5 List of sound engines

Creative/E-mu

Sound engine name	isng value	Creative/E-mu SF version	Bit depth	Sound cards
EMU8000	EMU8000	1.0, 1.5, 2.00	16 bit	AWE32, AWE64
EMU10K1	E-mu 10K1	2.01	16 bit	SB Live!
EMU10K2	E-mu 10K2	2.01	16 bit	SB Audigy
EMU20K1, EMU20K2	X-Fi	2.04	24 bit	SB X-Fi

SFe

Sound engine name	isng value	SFe version	Bit depth
SFe 4	SFe 4	4.0	32 bit

5.6.6 IC RD sub-chunk

To ease the creation of library management systems that are compatible with multiple languages, the naming convention for the IC RD sub-chunk has been changed.

The value of IC RD must now be compliant with the ISO-8601 standard. There are two valid formats:

- Date only: for example 2025-01-20
- Date and time: for example 2025-01-20T02:10:48Z

Library management systems should be able to read the value of the IC RD sub-chunk and show the date (and time if applicable) in the correct language in a field that can be sorted.

If the value of the IC RD sub-chunk is missing or not in any of the above two valid formats, the program may either:

- attempt to parse the value (if the chunk is present)
- ignore the value (if present) and show an "unknown" error on the date (and time) field
- overwrite the value with the current date (and time) if the program is an editor

The program must NOT reject a file with a missing or invalid IC RD sub-chunk as Structurally Unsound.

5.6.7 INAM, IENG, IPRD, ICOP, ICMT and ISFT sub-chunks

These sub-chunks are mostly the same as in legacy SF2.04, but UTF-8 is now used instead of ASCII, and the length limit is removed.

Reject anything not terminated with a zero byte. Do NOT reject the file as Structurally Unsound.

5.6.8 irom and iver sub-chunks

Read the legacy SF2.04 specification for info on how to use ROM samples.

The ROM emulator should be implemented in SFe programs.

5.6.9 SFty sub-chunk

The SFty sub-chunk is required and contains a case-sensitive UTF-8 string with even length identifying the type of format used in SFe. Its value is used by SFe-compatible players to assist in loading banks by telling the program what variant of SFe to load a bank as.

The defined values of the SFty chunk are:

- the 12 bytes representing SFe-static as 10 UTF-8 characters followed by two zero bytes.
- the 20 bytes representing SFe-static with TSC as 19 UTF-8 characters followed by one zero byte.
- the 20 bytes representing SFe-static (8-bit) as 18 UTF-8 characters followed by two zero bytes.
- the 28 bytes representing SFe-static (8-bit) with TSC as 27 UTF-8 characters followed by one zero byte.

The field should not be longer than 28 bytes in SFe 4.0.

If the SFty sub-chunk is missing or its contents are an undefined value or in an invalid format, other properties of the structure should be used to determine the variant of SFe that is in use. Do not assume SFe-static; only use such a value when it is evident beyond a reasonable doubt that the file used is in the SFe-static format.

5.6.10 SFvx sub-chunk

The SFvx sub-chunk is required and contains extended SFe version attributes. It is always 46 bytes in length, containing data in the structure below:

```
struct SFeExtendedVersion
{
    WORD wSFeSpecMajorVersion;
    WORD wSFeSpecMinorVersion;
    CHAR achSFeSpecType[20];
    WORD wSFeDraftMilestone;
    CHAR achSFeFullVersion[20];
};
```

The WORD values wSFeSpecMajorVersion and wSFeSpecMinorVersion contain the SFe specification version, and are used to differentiate between different SFe versions as the value of ifil only changes when the format of the `SFvx` sub-chunk does so.

The case-sensitive UTF-8 character field achSFeSpecType contains a specification type in UTF-8. For the purposes of this specification, the defined values are:

- Final for final specifications.
- Release Candidate for release candidate specifications.
- Milestone for draft specification milestones.
- Dev for rolling draft specifications.

Assume Final if contents are unknown.

The WORD value wSFeDraftMilestone contains the draft specification milestone or release candidate number that a bank was created to. This varies depending on the value of achSFeSpecType.

The case-sensitive UTF-8 character field achSFeFullVersion contains the full version string of the specification used, for example 4.0-rc2.

If the SFvx sub-chunk is missing or of an incorrect size, assume these values:

- wSFeSpecMajorVersion and wSFeSpecMinorVersion correspond to the highest version declared in the flag sub-chunk
 - If there is no valid flag sub-chunk, then assume the highest SFe version supported by the program.
- achSFeSpecType=Final
- wSFeDraftMilestone=0
- achSFeFullVersion corresponds to the other assumed values

The file may optionally be rejected as Structurally Unsound.

5.6.11 flag sub-chunk

The flag sub-chunk is required and contains the feature flags used by a bank. It is always a multiple of 6 bytes in length, and contains at least 2 records (1 feature flag and a record at the end) according to the structure:

```
struct SFeFeatureFlag
{
    BYTE byBranch;
    BYTE byLeaf;
    DWORD dwFlags;
};
```

The BYTE value `byBranch` represents the branch of the feature. Branches correspond to types of features.

The BYTE value `byLeaf` represents the leaf of the feature. Leaves correspond to specific features.

The DWORD value `dwFlags` represents the feature flags themselves, which represent different parts of the feature. Depending on the `byLeaf` value, it can be a number, a series of bytes, etc.

A tree value is a combination of a branch value and a leaf value, and is conventionally written in the format `branch : leaf` with hexadecimal values, for example "feature flag 03:01" refers to the feature flag with branch number 3 and leaf number 1 (SFe Compression sample compression formats). While the `flag` sub-chunk uses a tree structure, it should be noted that no branch includes sub-branches; the branches only include leaves.

Branch numbers between 240 (F0) and 255 (FF) are private-use branches that will not be defined in the SFe specification itself, and are free to be used by programs.

An exhaustive list of feature flags and their corresponding tree values can be found in section 6.2.

The final record should never be accessed in normal usage, but its value of `byBranch` and `byLeaf` have strict values depending on the specification version. Any records after the terminal record or with a higher tree value combination (except for the defined private-use area) should be ignored.

If the flag sub-chunk is missing or an incorrect size, then an effort should be made to recover the data. If data is not recoverable, then it can be rebuilt from the properties of the data in the rest of the bank. Do not reject the file as Structurally Unsound.

5.7 sdata-list chunk

5.7.1 smpl sub-chunk

This sub-chunk will now be present in most SFe files, as there is likely to be no ROM where samples can be read from. This does not include AWE ROM emulation. It works in an almost identical manner to legacy SF2.04, with these important differences:

- This contains one or more samples of audio in linearly coded 16-bit, signed words. These words are little-endian if the header is RIFF or RF64, and big-endian if the header is RIFX.
- No more leeway of 46 zero-valued samples is required after each sample.
- Before saving, SFe editors should insert this leeway. Otherwise, they might give a warning telling the user that loop and interpolation quality may be affected.
- If ROM samples are detected in SFe files, attempt to load them, even if this sub-chunk is missing.
- If this sub-chunk is missing, and no ROM samples are found, show a suitable error message.

5.7.2 SFe Compression

To implement compression in your SFe bank, please use the SFe Compression specification, listed in this section (5.7.2).

What is SFe Compression?

SFe Compression is the compression encoding system used by SFe, based on the earlier [Werner SF3](#) system widely used by the open source community.

By standardising on Werner SF3 in the form of SFe Compression, we will hopefully ensure that everyone uses the same compression formats. Due to this, we will only make small changes to SFe Compression which correspond to updates to the Werner SF3 specification by other SF player programs. To achieve this, all SFe players should implement SFe Compression.

File identification for SFe Compression

The `wMajor` value in the `ifil` sub-chunk is set to 3 instead of 2. The value of the `SFvx` sub-chunk remains unchanged. Therefore, SFe players should not use the `ifil` value to determine the SFe version, but rather the `SFvx` sub-chunk.

sfSampleType in shdr sub-chunk

Bit 4 of the `sfSampleType` field indicates a sample that has received some kind of compression. While most Werner SF3 compatible programs compress the samples using the Vorbis format, it cannot be assumed. SFe players must determine the encoding that is used for each sample.

For a sample to be valid:

- the type of compression and/or encoding must be recognised and supported by the program.
- the compression and/or encoding format must be valid.
- the compressed sample must only have one channel.

A sample is not valid if any of these conditions are not true. If a sample is not valid, then all instruments and presets that use the sample should be rejected.

Interpretation of sample data index fields in shdr sub-chunk

If bit 4 of the sfSampleType field is set, then the interpretation of the four sample data index fields changes:

- dwStart points to the first byte of the compressed byte stream relative to the beginning of the `smp1` sub-chunk.
- dwEnd points to the last byte of the compressed byte stream, instead of the first zero-valued sample data point after the sample data in legacy SF2.04.
- dwStartLoop and dwEndLoop specify loop points relative to the start of the individual uncompressed sample data, in sample data points.

If bit 4 of the sfSampleType field is clear, then the sample data index fields should be interpreted as in legacy SF2.04.

Using both compressed and uncompressed samples in the same file

You can use both compressed and uncompressed samples with SFe Compression. Simply place uncompressed PCM samples at the beginning of the `smp1` sub-chunk before the compressed sample byte stream. Because each sample is compressed individually, the resulting byte streams of all encoded samples are written to the `smp1` sub-chunk. The `smp1` sub-chunk may also contain uncompressed little-endian PCM samples.

For compressed byte streams, it is not necessary to add forty-six zero-valued sample data points after each sample. The length of the `smp1` sub-chunk is not required to be a multiple of two for compressed banks, and its surrounding LIST chunk is also not padded to a multiple of two as a consequence.

Unsupported features

Compressed samples are always 16-bit samples. They do not make use of the `sm24` (or `sm32`) sub-chunk. If an `sm24` sub-chunk is present, its respective byte counterparts to the compressed byte stream stored in the `smp1` sub-chunk remain unused. Since all uncompressed PCM samples are stored before compressed samples in the `smp1` sub-chunk, the size of the `sm24` sub-chunk is minimised. The `sm24` size constraint defined in legacy SF2.04 therefore no longer applies in compressed banks.

Sample links are not used in banks compressed with SFe Compression. The value of `wSampleLink` should be read and written as zero.

Proprietary compression formats

The only supported compression system for SFe is the Werner SF3-compatible SFe Compression. Proprietary SF compression formats (`.sfark`, `.sfpack`, `.sf2pack`, `.sfogg`, `.sfq`, `.sf4`) must not be used. Because Cognitone SF4-formatted banks are not valid Werner SF3 banks, they are also incompatible with SFe Compression.

5.7.3 sm24 and sm32 sub-chunk

These sub-chunks are optional.

- The sm32 sub-chunk contains the least significant byte, and the sm24 sub-chunk contains the next least significant byte after sm32.
- Each sub-chunk is exactly half the size of the smpl sub-chunk for uncompressed banks. This may not apply when SFe Compression is in use.
- For every two bytes in the smpl sub-chunk, there is one byte in these sub-chunks (if all samples are compressed).
- These sub-chunks can only be used with uncompressed samples. This limitation may be removed in future versions of SFe.

If these sub-chunks are present, they are combined with the other sub-chunks to create a sample with higher bitdepth.

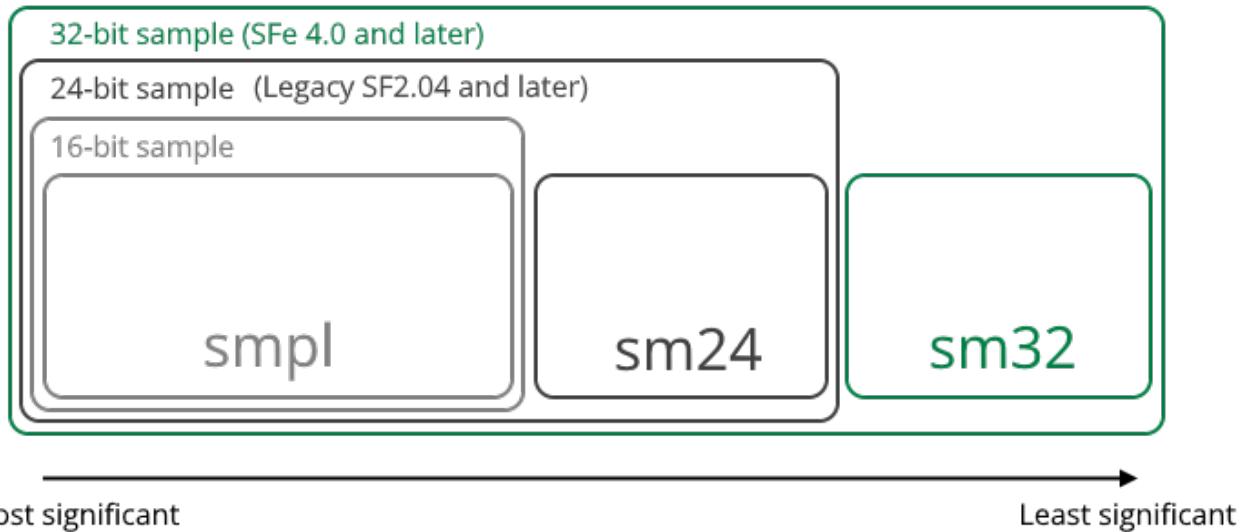


Figure 6: Available sample types in standard sample mode.

- If the `ifil` version is below 2.04 (signifying legacy SF2.01 or earlier), both `sm24` and `sm32` are ignored.
- If the `ifil` version is exactly 2.04 (signifying legacy SF2.04), only `sm32` is ignored. The `sm24` sub-chunk is still used.
- For uncompressed banks, these sub-chunks are not exactly half the size of the `smp1` sub-chunk (or otherwise invalid), the data should be ignored.
- If only the `sm32` sub-chunk is invalid, the `sm24` sub-chunk should still be loaded.
- However, if only the `sm24` sub-chunk is invalid, both sub-chunks should be ignored.
- If the `sm24` sub-chunk is ignored, the synthesizer should only attempt to render the first 16 bits of the samples contained within the `smp1` sub-chunk.
- If only the `sm32` sub-chunk is ignored, the synthesizer should attempt to render both the `smp1` and `sm24` sub-chunks, resulting in a 24-bit sample.
- We recommend only using `sm32` with 64-bit chunk headers. Using `sm32` with a 32-bit chunk header is syntactically valid, but it is not practical due to file size limitations

5.7.4 Using 8-bit samples

If the `smpl` sub-chunk is missing, but the `sm24` or `sm32` sub-chunks are present, then the `sm24` or `sm32` sub-chunk is considered "orphaned".

In this case, the value of the `SFty` sub-chunk inside the `ISFe-list` sub-chunk should be checked. If it is equal to `SFe-static` (8-bit) or `SFe-static` (8-bit) with `TSC`, then 8-bit sample mode is to be activated.

- The sample depth is eight bits.
- The length of the `sm24` sub-chunk should be rounded up to the nearest byte, as if the sub-chunk was being used with a `smpl` sub-chunk.
- This mode can only be used if all samples are 8-bit. You cannot mix 8-bit and higher sample depths.
- SFe 4 does not support non-standard sample bit depths (6-bit, 12-bit, 18-bit, etc.)

If there is an orphaned `sm24` or `sm32` sub-chunk, and the `SFty` sub-chunk is missing or is not equal to a value that corresponds to 8-bit samples being present, then the file should be rejected as Structurally Unsound.

If there is both an orphaned `sm24` and an orphaned `sm32` sub-chunk, the `sm24` sub-chunk is the most significant byte, and 16-bit sample playback is assumed. Editing software should give a warning if there is a `sm24` and a `sm32` sub-chunk but no `smpl` sub-chunk, and should convert it to a 2.01-compliant 16-bit format. This behaviour should be followed if supported, regardless of whether 8-bit sample playback support is actually supported.

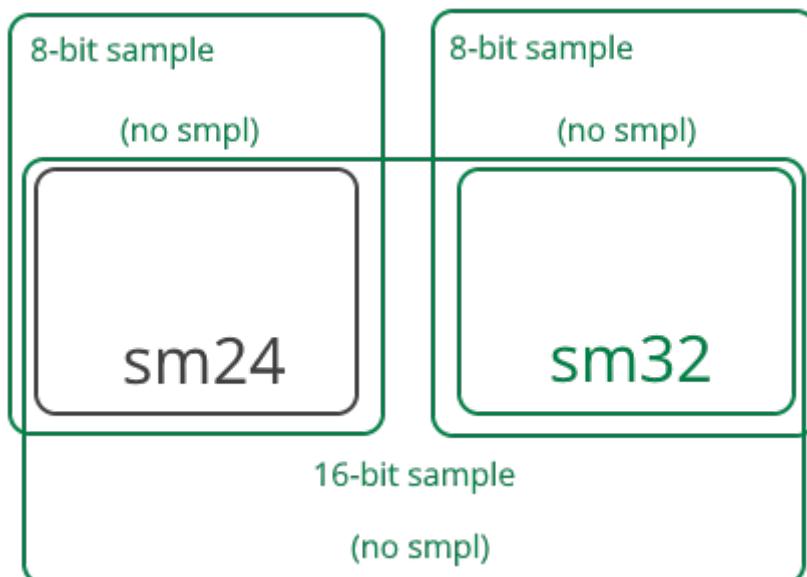


Figure 7: Available sample types in 8-bit sample mode.

5.7.5 Looping rules

- No more leeway of eight samples is required.
- Before saving, SFe editors might give a warning about this leeway telling the user that loop and interpolation quality may be affected.

5.8 pdta-list chunk

5.8.1 phdr sub-chunk

Its size is a multiple of 38 bytes, and its structure is the same as in legacy SF2.04.

The last sfPresetHeader entry shouldn't need to be accessed, apart from the uses described in SFSPEC24.PDF. The phdr sub-chunk is required; files without a phdr sub-chunk are Structurally Unsound.

5.8.2 New Bank System

In SFe 4.0, the bank system has been completely overhauled. Please read this section carefully to ensure that you correctly implement bank selects in your program.

Using MIDI Control Change #32 (Bank Select LSB)

In legacy SF2.04, the wBank field stores the bank that the preset can be found in. Due to a forward-thinking decision by E-mu, it is a WORD (16-bit) instead of a BYTE (8-bit). This means that it could theoretically store values for both bank select instructions found in MIDI 1.0.

Bank select LSB support is added by using the unused 8 bits of wBank according to the figure below. Bits 2-8 of *both* byBankMSB and byBankLSB are now used to set a bank change.

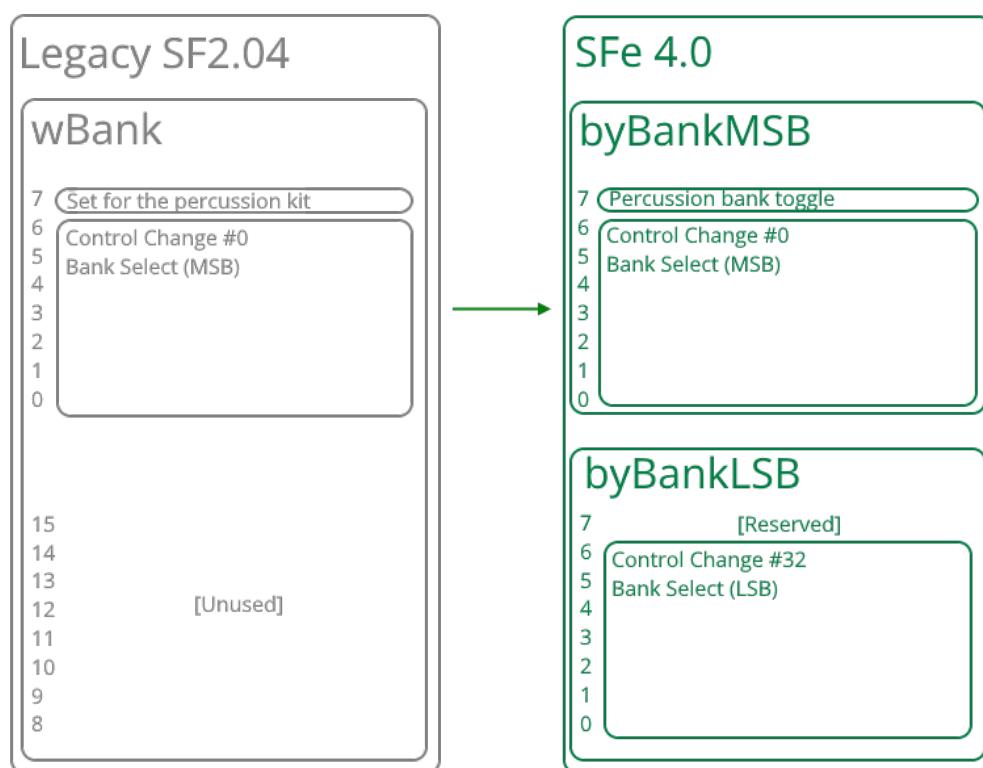


Figure 8: How the bank select logic differs from legacy SF2.04.

Introducing byBankMSB and byBankLSB

In the above figure, wBank has been replaced with byBankMSB and byBankLSB.

This splits the one WORD in legacy SF2.04 into two BYTE values, one for each bank. byBankMSB goes before byBankLSB due to RIFF being a little-endian format. For RIFX, byBankLSB is first.

Using more than one percussion bank

Legacy SF2.04 allows bank developers to define one bank of percussion kits for use in channel 10 that can be switched between using MIDI Program Change instructions by using the wBank number 128. In other words, if bit 7 is set, bits 0-6 must be clear - you cannot use bank select instructions with channel 10.

SFe 4.0 now allows users to set bit 7 with any value for bits 0-6. The result is that there are 128 percussion banks available when using byBankMSB, as shown by the figure below.

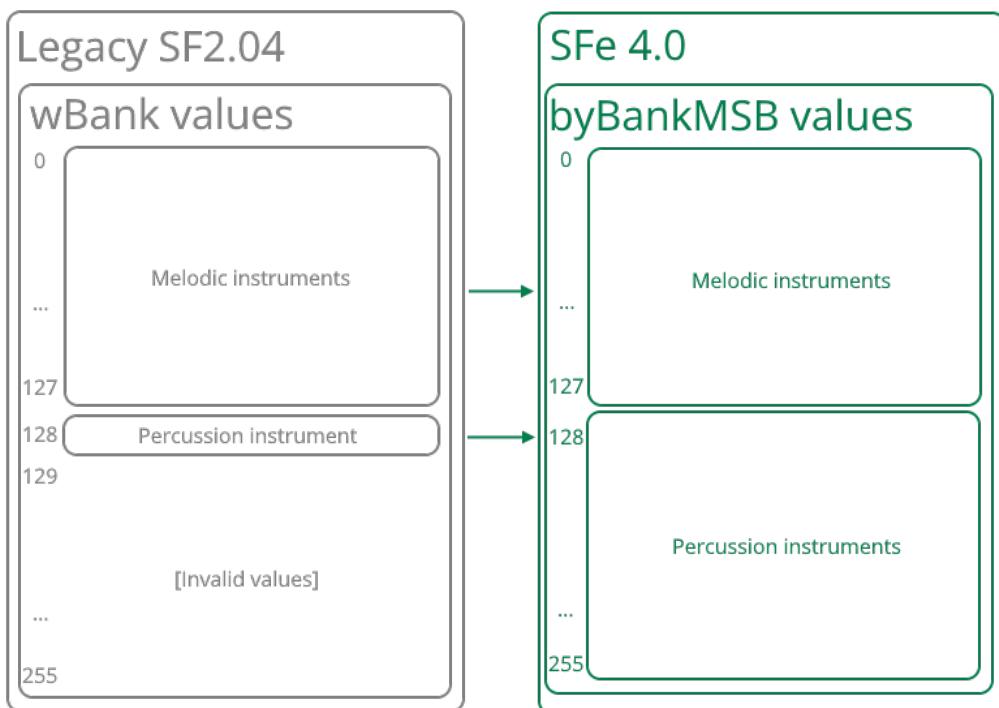


Figure 9: How the percussion bank listing differs from legacy SF2.04. When byte 7 is set for byBankMSB, byBankLSB may also be used. Therefore, a total of 16384 (128×128) banks of percussion kits may be used.

Flowchart for correct handling of bank select instructions

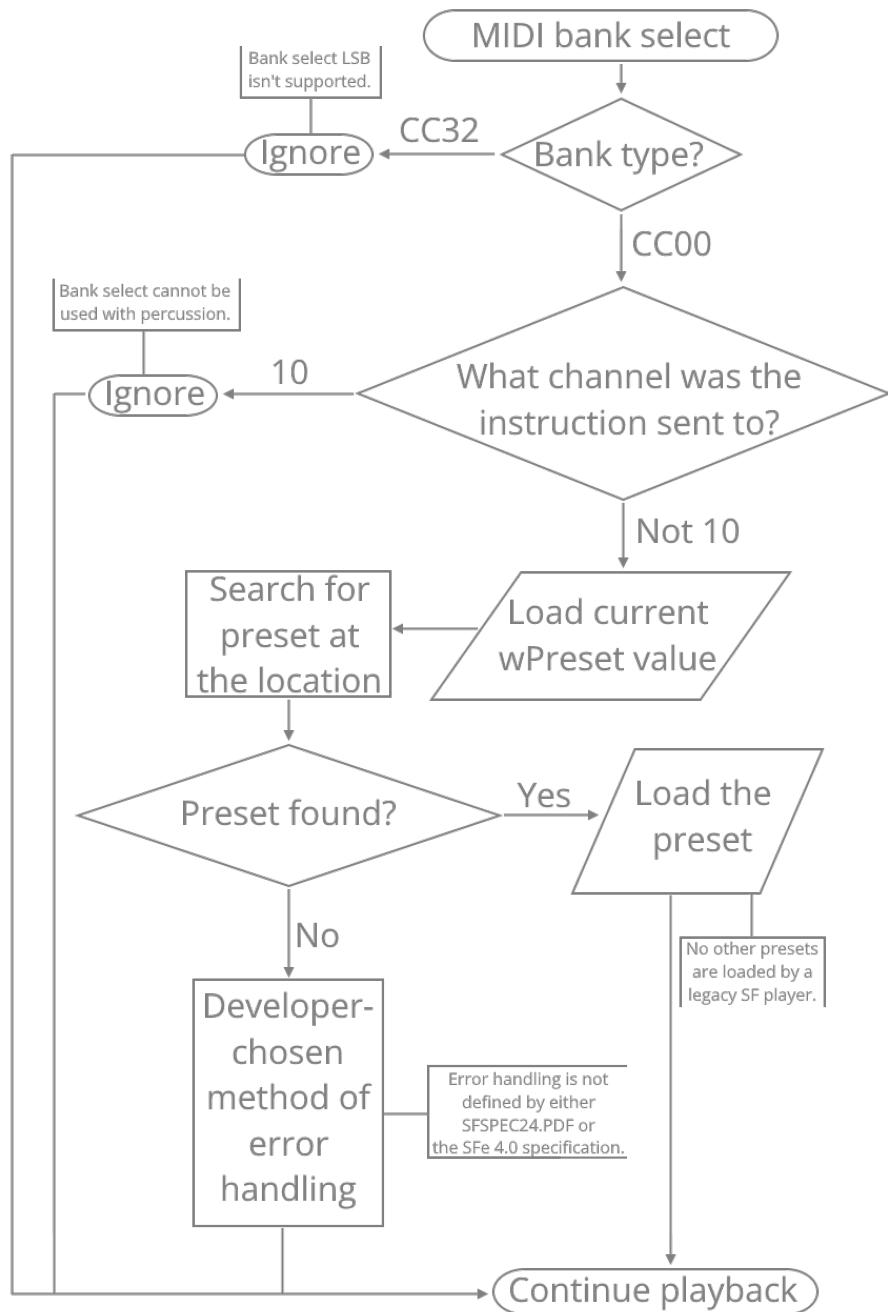


Figure 10: The flowchart for bank select instructions in legacy SF2.04.

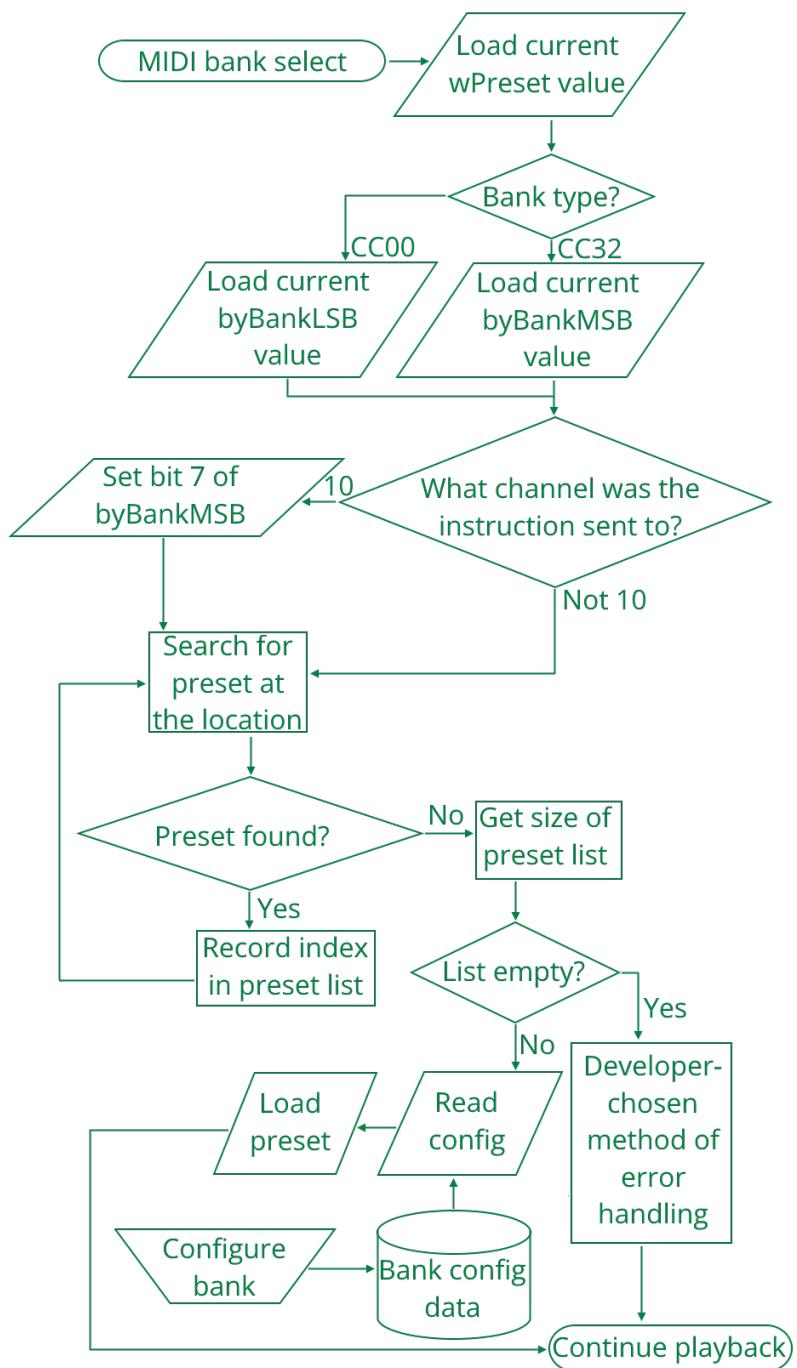


Figure 11: The flowchart for bank select instructions in SFe 4.0.

Notice that not only are extra steps added for bank select LSB and percussion bank select handling, but extra configuration information used by the player is added to determine the correct preset to use.

5.8.3 inst sub-chunk

Its size is a multiple of 22 bytes, and its structure is the same as in legacy SF2.04.

The `inst` sub-chunk is required; files without an `inst` sub-chunk are Structurally Unsound.

5.8.4 shdr sub-chunk

Its size is a multiple of 46 bytes, and its structure is the same as in legacy SF2.04.

The `shdr` sub-chunk is required; files without a `shdr` sub-chunk are Structurally Unsound.

Sample Rate Limit Changes

- In SFe, sample rates (`dwSampleRate`) are stored as a 32-bit integer. This is the same behavior as seen in the legacy SF2.04 format. This results in a theoretical maximum sample rate of 4,294,967,295 Hz.
- In the legacy SF2.04 specification, E-mu suggested that sample rates of below 400 Hz or above 50,000 Hz should be avoided as some legacy hardware platforms may not be able to reproduce these sounds. This is not a limitation of the specification, but rather a limitation of legacy sound cards.
- Despite this, Creative did not use 16-bit integers for sample rate in legacy SF2.04. It is thus safe to use sample rates in excess of 50,000 Hz. If a sample rate of below 400 Hz or above 50,000 Hz is encountered, no attempt should be made to change the sample rate.
- A zero sample rate should be reset.

sfSampleType and SFe Compression

Bit 4 of `sfSampleType` is reserved for SFe Compression usage.

- Read section 6.2 for more information on SFe Compression!

5.8.5 Other sub-chunks

The `pbag`, `pmod`, `pgen`, `ibag`, `imod` and `igen` sub-chunks work in the same way as legacy SF2.04. Read `SFSPEC24.PDF` for more information.

If any of the sub-chunks listed above is missing or invalid, the SFe bank is Structurally Unsound.

Section 6

SFe enumerations and feature flags

6.1 About SFe enumerations

SFe 4.0 enumerations are identical to legacy SF2.04, but more enumerations may be defined in the future.

Once the enumeration model is changed significantly, we will list all changed modulator and generator enums with a clear explanation of what they do. Equations for timecents will also be included!

6.2 Feature flags

6.2.1 Feature flag tree structure

The feature flags system is split like this:

- Branches: Roughly corresponds to each version (but not all listed features are part of the version). Maximum of 256. Number may increase with later `wMajor` versions.
- Leaves: Corresponds to each feature. Maximum of 256. These may change with later `wMajor` versions. Contains 32-bit data declaring how much of the feature is implemented.
- Flags: Each of the 32 bits that comprise a leaf, declaring support for specific features.

Feature flags listed as "reserved" must not be used for private use. Branches 240 (F0) to 255 (FF) are provided for such use.

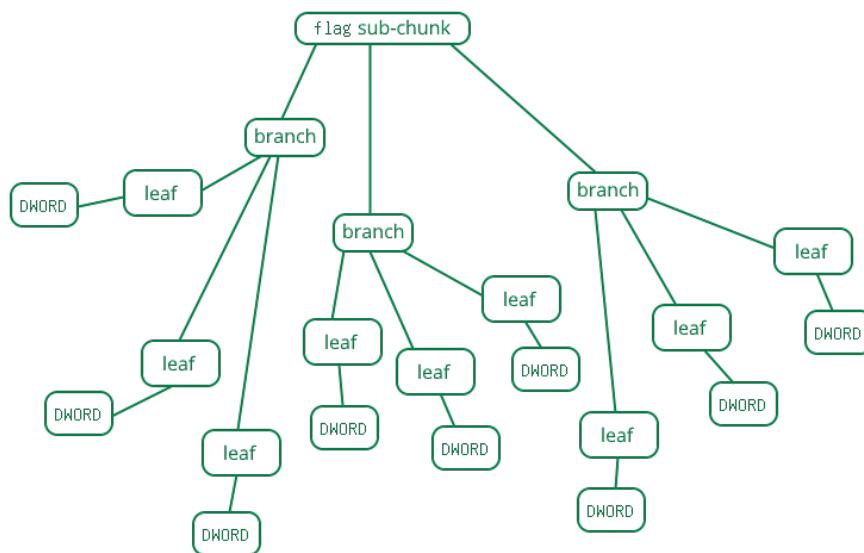


Figure 12: The tree structure of the feature flags system.

6.2.2 Branch 00 Foundational synthesis engine

00:00 Tuning

- Bit 1: Coarse tuning
- Bit 2: Fine tuning
- Bit 3: Root key
- Bit 4: Scale tuning

00:01 Looping

- Bit 1: Loop during release
- Bit 2: Non-loop during release

00:02 Filter Types

- Bit 1: Sound Blaster compatible low pass (12dB filter)

00:03 Filter Parameters

- Bits 1-16: Maximum supported filter frequency
- Bits 17-24: Maximum supported filter resonance
- Bits 25-32: Reserved

00:04 Amplification and attenuation

- Bit 1: Attenuation supported in preset (0.4x)
- Bit 2: Attenuation supported in instrument (0.4x)
- Bit 3: Amplification supported in preset (0.4x)
- Bit 4: Reserved
- Bit 5: Reserved
- Bit 6: Reserved
- Bit 7: Reserved

00:05 Effects blocks

- Bit 1: Instrument-level reverb
- Bit 2: CC91 reverb
- Bit 3: Combined reverb
- Bit 4: Adjustable reverb
- Bit 9: Instrument-level chorus
- Bit 10: CC93 chorus
- Bit 11: Combined chorus
- Bit 12: Adjustable chorus
- Bit 17: Pan supported

00:06 Low Frequency Oscillators

- Bit 1: Vibrato supported
- Bit 2: Pitch Modulation
- Bit 3: Filter Modulation
- Bit 4: Amplitude Modulation

00:07 Envelopes

- Bit 1: Volume delay
- Bit 2: Volume attack
- Bit 3: Volume hold
- Bit 4: Volume decay
- Bit 5: Volume sustain
- Bit 6: Volume release
- Bit 7: Key to volume hold
- Bit 8: Key to volume decay
- Bit 9: Modulation delay
- Bit 10: Modulation attack
- Bit 11: Modulation hold
- Bit 12: Modulation decay
- Bit 13: Modulation sustain
- Bit 14: Modulation release
- Bit 15: Key to modulation hold
- Bit 16: Key to modulation decay
- Bit 17: Modulation of volume
- Bit 18: Modulation of pitch
- Bit 19: Modulation of filter

00:08 MIDI Control Changes

- Bit 1: 00 Bank Select MSB
- Bit 2: 00 Bank Select MSB for percussion
- Bit 3: 06 Data Entry MSB
- Bit 4: 32 Bank Select LSB (Multiple banks)
- Bit 5: 32 Bank Select LSB (Preset name)
- Bit 6: 32 Bank Select LSB (byBankLSB support)
- Bit 7: 32 Bank Select LSB for percussion
- Bit 8: 38 Data Entry LSB
- Bit 9: 64 Sustain
- Bit 10: 66 Soft
- Bit 11: 67 Sostenuto
- Bit 12: 98 NRPN LSB
- Bit 13: 99 NRPN MSB
- Bit 14: 100 RPN LSB
- Bit 15: 101 RPN MSB
- Bit 16: 120 All sound off
- Bit 17: 121 Reset all controllers
- Bit 18: 123 All notes off
- Bit 19: Reserved

00:09 Generators

- Bit 1: Index gen support
- Bit 2: Range gen support
- Bit 3: Substitution gen support
- Bit 4: Sample gen support
- Bit 5: Value gen support
- Bit 6: PGEN support
- Bit 7: IGEN support

00:0a Zones

- Bit 1: Key range
- Bit 2: Velocity range
- Bit 3: Exclusive class
- Bit 4: Fixed key
- Bit 5: Fixed velocity
- Bit 6: Sample offset
- Bit 7: Loop offset

00:0b Reserved

- Bit 1: Reserved
- Bit 2: Reserved
- Bit 3: Reserved
- Bit 4: Reserved

6.2.3 Branch 01 Modulators and NRPN

01:00 Modulators

- Bit 1: Primary source
- Bit 2: Secondary source
- Bit 3: Transform support
- Bit 4: Linear curves
- Bit 5: Concave curves
- Bit 6: Convex curves
- Bit 7: Switch curves
- Bit 8: Positive unipolar
- Bit 9: Negative unipolar
- Bit 10: Positive bipolar
- Bit 11: Negative bipolar
- Bit 12: Modulator chaining
- Bit 13: PMOD support
- Bit 14: IMOD support

01:01 Modulation controllers

- Bit 1: Note-on velocity
- Bit 2: Note-on key number
- Bit 3: Poly pressure
- Bit 4: Channel pressure
- Bit 5: Pitch wheel
- Bit 6: Pitch wheel sensitivity

01:02 Modulation parameters 1

- Bits 1-32: MIDI CC000-031

01:03 Modulation parameters 2

- Bits 1-32: MIDI CC064-095

01:04 Modulation parameters 3

- Bits 1-32: MIDI CC096-127

01:05 NRPN

- Bit 1: NRPN select MSB=120
- Bit 2: NRPN select LSB: 1-2 digits
- Bit 3: NRPN select LSB: 3 digits
- Bit 4: NRPN select LSB: 4 digits
- Bit 5: NRPN select LSB: 5 digits

01:06 Default modulators

- Bit 1: MIDI note on to initial attenuation
- Bit 2: MIDI note on to filter cutoff
- Bit 3: MIDI channel pressure to vibrato LFO pitch depth
- Bit 4: MIDI CC1 to vibrato LFO pitch depth
- Bit 5: MIDI CC7 to initial attenuation
- Bit 6: MIDI CC10 to pan position
- Bit 7: MIDI CC11 to initial attenuation
- Bit 8: MIDI CC91 to reverb send
- Bit 9: MIDI CC93 to chorus send
- Bit 10: MIDI pitch wheel to initial pitch, controlled by pitch wheel sensitivity
- Bit 17: MIDI note on to filter cutoff (SF2.00)
- Bit 18: MIDI note on to filter cutoff (SF2.01)
- Bit 19: MIDI note on to filter cutoff (SF2.04)
- Bit 24 off, bit 25 off: Reserved
- Bit 24 on, bit 25 off: Reserved
- Bit 24 on, bit 25 on: Reserved

01:07 Reserved

- Bit 1 off, bit 2 off: Reserved
- Bit 1 on, bit 2 off: Reserved
- Bit 1 on, bit 2 on: Reserved
- Bit 3: Reserved
- Bit 4: Reserved
- Bit 5: Reserved
- Bit 6: Reserved

6.2.4 Branch 02 Sample bitdepth support

02:00 24-bit support

- Bit 1 off, bit 2 off: No support
- Bit 1 on, bit 2 off: Read support only
- Bit 1 on, bit 2 on: Playback support
- Bit 3: 8-bit support

02:01 32-bit support

- Bit 1 off, bit 2 off: No support
- Bit 1 on, bit 2 off: Read support only
- Bit 1 on, bit 2 on: Playback support
- Bit 2: Support for combining two 8-bit chunks into a 16-bit sample

6.2.5 Branch 03 SFe Compression support

03:00 Compression flag

- 0: sfSampleType bit 4 unsupported
- 1: sfSampleType bit 4 supported

03:01 Sample compression formats

- Bit 1: OGG
- Bit 2: Opus
- Bit 3: FLAC
- Bit 4: MP3

6.2.6 Branch 04 Metadata upgrades

04:00 Metadata improvements

- Bit 1: UTF-8 in INFO
- Bit 2: UTF-8 in pdta

04:01 Reserved

- Bit 1: Reserved
- Bit 2: Reserved
- Bit 3: Reserved

04:02 User-defined sample ROMs

- Bit 1: Support for user-defined sample ROMs

04:03 ROM emulator

- Bit 1: 1MB ROM emulator support

Section 7

Parameters and synthesis model

7.1 About the synthesis model

Currently, the synthesis model of SFe 4.0 is identical to that of legacy SF 2.04. However, more changes are coming soon.

When the synthesis model is changed, there will be a detailed description of everything here.

7.2 MIDI functions

7.2.1 MIDI bank select

Control Change #32 (Bank Select "LSB") has been modified to use the `byBankLSB` value.

7.2.2 Other MIDI functions

All other MIDI functions are unchanged from legacy SF2.04.

7.3 Parameter units, generators, modulators and NRPNs

All of these are the same as legacy SF2.04.

7.4 On implementation accuracy

E-mu was very lax when it came to accuracy of legacy SF implementations. This was because of the limitations of computers and legacy soundcards that the legacy SF spec and its implementations were initially designed to run on. While this was the only way that legacy SF could gain the popularity that it did with software implementations, this meant that bank developers often had to declare the program that their file was intended to be used with. This hampered interoperability with different SF players, including those that may have been embedded into musical instruments.

Because today's computers are much faster, and legacy soundcards are no longer in widespread use, the requirements for implementation accuracy in SFe are far more strict. All SFe players are required to recognise the `f`lag sub-chunk (section 5.12.2) and warn the user if there is a mismatch between feature flags in the bank and the program's support.

While the `f`lag sub-chunk will be useful to alert users about incompatible programs, program developers should make an effort to ensure that their program is 100% compatible.

Section 8

SFe error handling

8.1 Structurally Unsound errors

"Structurally Unsound" errors are those defined by E-mu (in legacy SF2.04), Werner Schweer (in Werner SF3) and the SFe Team (in SFe) to prevent the bank from working properly in a way that means that it can not be used. These errors must be fixed before an SF player can load it, unless the SF player implements SFe automatic repair.

The error correction process for structural errors in SFe is slightly different from that in legacy SF2.04:

- If a ds64 chunk, BW64 or RF64 header is found in a file, SFe players that do not support 64-bit chunk headers should output a specific error, as mentioned in the SFe program specification.
- If the ds64 chunk is missing, the bank uses a 32-bit chunk header. Make sure that ckSize is accurate, using the same techniques as for SF2.04.
 - If the value of ckSize is 4,294,967,295 (like in ds64), or any other inaccurate value, reject the file as "Structurally Unsound."
 - If the value of ckSize is correct, open the file and output a warning message, as mentioned in the SFe program specification.
 - More advanced programs may also recognize files larger than 4GiB with 32-bit headers. However, such a file should be repaired with a 64-bit header as soon as possible.
- Banks with 64-bit headers, but with a ckSize value that is not 4,294,967,295 should be rejected as "Structurally Unsound," as this is not valid in RIFF64.

8.2 Non-critical errors

"Non-critical" errors are errors that mean that data cannot be read properly, however as they are not Structurally Unsound errors, the file can be loaded. What usually happens is that the damaged data is ignored, and the rest of the bank is loaded.

While non-critical errors don't prevent the use of the bank, it is important that they are fixed to ensure that the bank functions work as intended on all SF players that conform to a specification, whether that be legacy SF2.04, Werner SF3 or SFe.

8.3 Duplicated preset locations within files

This occurs when the file is structurally damaged or manually edited in a manner where more than one preset has the same value of `byBankMSB`, `byBankLSB` and `wPreset` (for instance, `015:000:081`).

- In legacy SF2.04, the first preset in the location would be used by default, but the other presets would still be retained.
- Such other presets must be moved before they can be used.
- In SFe, the preset to be used is no longer necessarily the first one. Instead, selecting the correct preset (or combination of presets) to use will be permissible.
- Such a feature is optional, and if not implemented, the player should use the legacy SF2.04 behavior in these cases (use the first preset found).
- Editors should warn the user if such presets are found.
- This behavior might change in future versions, so please take the `ifil` value, and later versions of this specification, into account.

8.4 Duplicated preset locations across files

This occurs when multiple files are loaded simultaneously (now a required feature for SFe), but some or all of the files have presets with identical byBankMSB, byBankLSB and wPreset values.

- Behavior was undefined in legacy SF2.04.
- This was because multiple file loading was not a standard feature mandated in the legacy SoundFont® standard.
- Legacy SF2.04 and Werner SF3 compatible software developers therefore had the liberty to implement multiple file loading; however, they wanted to.
- This edge case will now be defined in SFe.
- If multiple presets across loaded files have the same value of byBankMSB, byBankLSB and wPreset, then the preset to be used may be selectable from all the presets with the same bank location (in the way described in section 8.3).
- Such a feature is also optional, and if not implemented, the player should use the legacy SF2.04 behavior in these cases (use the first preset found).
- Editors should warn the user if such presets are found.
- This behavior might change in future versions, so please take the ifil value, and later versions of this specification, into account.

8.5 Undefined chunks

Legacy SF2.04 players should ignore SFe-specific sub-chunks, as prescribed by E-mu.

Also, SFe 4.0 compatible players, which do not support future versions, should ignore sub-chunks which are used in future versions.

8.6 Unknown Enumerators

Any SFvx versions of 4.1 or above may have unknown enumeration values for an SFe 4.0 player.

This is entirely expected, and if unknown enumeration values are found, the Generator/Modulator should be ignored.

8.7 Maximum File Size Limit Exceeded

This occurs when the user loads a file that is larger than the maximum size that the SFe program can accommodate.

- In the AWE32 and AWE64, the limit was dependent on the memory installed on the card, which was up to 28 MiB.
- Later, the sound cards allowed the user to load files with file sizes up to system memory.
- SFe has defined limits that are found in the separate program specification.
- If these limits are reached, you can reject loading of the file with an error, or attempt to load the file anyway.
- SFe programs should warn the user if processing was automatically done to a file to reduce the file size to be in range.
- If multiple files are loaded, and the limit is reached, the order of files to be loaded can be defined by the author of the SFe-compatible software.

8.8 MIDI Errors

If a non-existent bank/preset combination is selected, the software should not playback any sound.

8.9 Illegal parameter values, out of range values, missing required items and illegal enumerators

These are handled as in legacy SF2.04.

Section 9

SiliconSFe

9.1 SiliconSFe overview

While we are unaware of any shipping products using the SiliconSF system found in SFSPEC24.PDF (the AWE cards used an early predecessor of SiliconSF), you can use ROM samples formatted in the SiliconSF format with SFe.

9.2 Header format

9.2.1 About the header format

The SiliconSFe header format is almost identical to legacy SF2.04, however an explanation is provided here due to poor documentation of SiliconSF.

Here is the SiliconSFe header format:

```
typedef struct romHdrType{
    DWORD romRiffHdr;
    DWORD romByteSize;
    CHAR romInterleaveIndex;
    CHAR romRevision[3];
    CHAR romVer[4];
    SHORT bankChecksum;
    SHORT bankChecksum2sComplement;
    CHAR bankSFeVersion;
    CHAR bankProduct[16];
    BYTE bankSampleCompType;
    CHAR filler1[2];
    CHAR bankStyle[16];
    CHAR bankCopyright[80];
    DWORD romSFeBankStart;
    DWORD romSineWaveStart;
    DWORD filler2[124];
    SHORT sampleSineWave[SINEWAVESIZE];
} romHdr;
```

9.2.2 romRiffHeader

In SiliconSFe, it is defined as the FourCC used by the chunk header type used by the integrated SF bank, for example RIFF, RF64, RIFD, etc.

In the legacy SF2.04 specification, this is named `romRsrc` and was declared by Creative as "unused". The name in SiliconSFe more accurately describes its usage.

9.2.3 romByteSize

This is an UNSIGNED DWORD value with the size of the SiliconSFe ROM blob in bytes. It is limited to 4 GiB in SiliconSFe 1.0. Signed integers are prohibited.

9.2.4 romInterleaveIndex

This is used for interleaved ROMs. You can interleave up to 256 ROMs with one SiliconSFe blob.

In the legacy SF2.04 specification, this is named `interleaveIndex`.

9.2.5 romRevision

This is a revision identifier as an integer. It is 3 bytes long and is independent from the version number found in `romVer`.

In the legacy SF2.04 specification, this is named `revision`.

9.2.6 romVer

This corresponds to the `iver` value in the integrated SF bank.

In the legacy SF2.04 specification, this is named `id` and is erroneously listed as corresponding to the `irom` value. The name in SiliconSFe more accurately describes its usage.

9.2.7 bankChecksum

This stores the CRC-16 (ARC) checksum of the integrated SF bank.

In the legacy SF2.04 specification, this is named `checksum`.

9.2.8 bankChecksum2sComplement

This stores the twos-complement of the value found in `bankChecksum`.

In the legacy SF2.04 specification, this is named `checksum2sComplement`.

9.2.9 bankSFeVersion

This value should be the same as the `wSFeSpecMajorVersion` value in the `SFvx` sub-chunk in SFe, and the same as the `wMajor` value in the `ifil` sub-chunk in non-SFe. For an unknown or other format, this value is 0.

In the legacy SF2.04 specification, this is named `bankFormat` and was declared by Creative as "unused".

9.2.10 bankProduct

This is a UTF-8 string that stores the product name (conventionally SiliconSFe).

In the legacy SF2.04 specification, this is named `product`.

9.2.11 sampleCompType

For the purpose of SiliconSFe, this value is 1 if any kind of sample precompensation is used, and 0 otherwise.

In the legacy SF2.04 specification, Creative said that it indicates the type of sample precompensation that is used in the SiliconSF blob.

9.2.12 bankStyle

This is a UTF-8 string that describes the musical style of the contents of the integrated SF bank.

In the legacy SF2.04 specification, this is named `style`.

9.2.13 bankCopyright

This is a UTF-8 string that stores copyright information about the SiliconSFe blob.

In the legacy SF2.04 specification, this is named `copyright`.

9.2.14 romSFeBankStart

This stores the location in the SiliconSFe blob where the integrated SF bank starts.

In the legacy SF2.04 specification, this is named `sampleStart`. The name in SiliconSFe more accurately describes its usage.

9.2.15 romSineWaveStart

This stores the location in the SiliconSFe blob where the test sine wave sample starts.

In the legacy SF2.04 specification, this is named `sineWaveStart`.

9.2.16 sampleSineWave

This contains SHORT values that correspond to a sine wave sample.

In the legacy SF2.04 specification, this is named `sineWave`.

9.3 AWE ROM emulator

9.3.1 Introducing the AWE ROM emulator

Except for when size concerns prohibit its inclusion, SFe players should include an AWE ROM emulator.

- The AWE ROM emulator includes 152 samples.
- The file size should be 1MB, as all samples should be to the same standard as the original.
- Samples which the program developer has the right to use can be used as a replacement for the original ROM samples.
- Freelyusable reference implementation samples are available, but are not intended for production use.
- Sample names will remain the same, but there will be acceptable alias names.
- Emulators should also be able to open up SF files (either legacy SF or SFe) containing samples and metadata.
- There may or may not be instruments or presets.

9.3.2 ROM emulator sample specification

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
1	acbasse1	Acoustic Bass E1	1535	1408-1530	
2	accordax2	Accordion A#2 Accordian A#2	620	566-616	Either spelling is acceptable
3	accordfx2	Accordion F#2 Accordian F#2	1049	979-1044	Either spelling is acceptable
4	accordfx3	Accordion F#3 Accordian F#3	858	794-854	Either spelling is acceptable
5	acgtrb3	Acoustic Guitar B3	6241	6168-6236	
6	acgtrg2	Acoustic Guitar G2	4882	4800-4877	
7	agogolotone	Agogo Low Tone Agogo Bell	4467	7-4463	
8	applause	Applause	8161	7-8156	
9	arcocelloax2	Arco Cello A#2	847	799-842	
10	arcocelld2	Arco Cello D2	1200	1127-1195	
11	arcoviolinc3	Arco Violin C3	1767	1702-1762	
12	arcoviolinc4	Arco Violin C4	1634	1569-1629	
13	arcovioline3	Arco Violin E3	1086	1035-1081	
14	arcoviolingx2	Arco Violin G#2	1732	1691-1727	
15	arcoviolingx3	Arco Violin G#3	1075	1032-1070	
16	asaxc2	Alto Sax C2	1150	1054-1145	
17	asaxc4	Alto Sax C4	1191	1130-1186	
18	asaxd3	Alto Sax D3	696	643-691	
19	asaxe2	Alto Sax E2	1228	1150-1223	
20	asaxf3	Alto Sax F3	910	863-905	
21	asaxg2	Alto Sax G2	1639	1567-1634	

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
22	bagpipedrna	Bag Drone A1 Bag Drone	1921	1764-1913	Based on name in 8MB E-mu bank
23	banjod3	Banjo D3	3583	3540-3579	
24	banjog2	Banjo G2	2850	2784-2845	
25	bassguitloop	Bass Guitar Loop	125	9-120	
26	bassoonc2	Bassoon C2 Bassoon	1059	938-1053	Only sample of that instrument
27	bd15	Bass Drum Kick Drum	1603	7-1599	
28	belltree	Bell Tree	7466	6263-7461	
29	bockclave	Claves	1774	7-1770	
30	brasssectc3	Brass Section C3	5600	5533-5595	
31	brasssectf5	Brass Section F5	5581	4989-5035	
32	bsawtoothwavea3	Sawtooth Wave Sample Bass SawtoothWave A3 Bass Sawtooth Wave	70	15-65	Real time synthesis activated without "Sample" suffix
33	cabasastrk	Cabasa	2679	7-2675	Based on name in 8MB E-mu bank
34	chanterax1	Bagpipe A#1 Bagpipe Chanter A#1 Chanter	1858	1802-1850	Based on name in 8MB E-mu bank Only sample of that instrument
35	chcrash	China Crash Cymbal	9700	6162-9695	Based on name in 8MB E-mu bank
36	clarinetb2	Clarinet B2	701	586-695	
37	clarinetd2	Clarinet D2	677	596-671	
38	clavc2	Clav C2 Clav	2985	2836-2980	Only sample of that instrument
39	coldglass7wave	Tinkle Bell Wave 1 Tinker Bell Wave 1	70	17-65	The alias is because of its use in the 1MB E-mu bank. Either spelling is acceptable

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
40	coldglass12wave	Tinkle Bell Wave 2 Tinker Bell Wave 2	91	15-86	The alias is because of its use in the 1MB E-mu bank. Either spelling is acceptable
41	contraviobass	Contrabass	1443	1302-1438	
42	cowbell	Cowbell	1760	7-1756	
43	crash5	Crash Cymbal	13534	8024-13529	
44	distgtra2	Distortion Guitar A2 Dist Gtr A2	1832	1745-1827	Based on name in 8MB E-mu bank
45	distgtra3	Distortion Guitar A3 Dist Gtr A3	1243	1195-1238	Based on name in 8MB E-mu bank
46	distgtrd4	Distortion Guitar A4 Dist Gtr A4	1766	1593-1761	Based on name in 8MB E-mu bank
47	distgtre3	Distortion Guitar E3 Dist Gtr E3	1432	1372-1427	Based on name in 8MB E-mu bank
48	dlcimrc3	Dulcimer C3 Dulcimer	3835	3778-3830	Only sample of that instrument
49	ebongostone	H Bongo Rim	3204	7-3200	Based on name in 8MB E-mu bank
50	elguitard2	Electric Guitar D2 Electric Guitar	1481	1401-1476	Only sample of that instrument
51	enghorndx3	English Horn D#3 English Horn	1540	1479-1534	Only sample of that instrument
52	epiano2ms	E Piano 2	1173	1120-1168	Only sample of that instrument
53	femalevoiceg2	Female Voice G2 Female Voice	8759	338-8755	Only sample of that instrument
54	filtersnap	Filter Snap	420	7-414	
55	floortombrite	Floor Tom Acoustic Tom	7172	5236-7167	Based on name in 8MB E-mu bank
56	flutec4	Flute C4	1432	1365-1427	Only sample of that instrument
57	frenchhorn4	French Horn G4 French Horn	1485	1420-1480	Only sample of that instrument
58	Fretlessa2	Fretless Bass A2 Fretless Bass	2341	2165-2336	Only sample of that instrument

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
59	glockloopc4	Glockenspiel C4 Glockenspiel Glockenspiel Loop	216	7-211	Only sample of that instrument
60	gsbassd2	Finger Bass D2 Finger Bass	905	750-900	Only sample of that instrument
61	guiro2	Guilo Up	2764	7-2759	Based on name in 8MB E-mu bank
62	guirodown	Guilo Down	2648	7-2644	
63	guitar1	Bandoneon	140	69-135	The alias is because of its use in the 1MB E-mu bank.
64	guitarfret	Guitar Fret Fret Noise	3572	7-3567	Based on name in 8MB E-mu bank
65	gunshot	Gun Shot	5396	7-5392	Based on name in 8MB E-mu bank
66	harmguitard3	Guitar Harmonics D3 Guitar Harmonics GtrHarmonics D3	344	298-338	Based on name in 8MB E-mu bank Only sample of that instrument
67	harmonicaa3	Harmonica A3 Harmonica	974	915-969	Only sample of that instrument
68	harpsichordc3	Harpsichord C3 Harpsichord	1391	1294-1386	Only sample of that instrument
69	hatopenms	Open High Hat	11710	5828-11705	Based on name in 8MB E-mu bank
70	hrmnmutec3	Harmon Mute C3	1485	1420-1481	
71	hrmnmutec4	Harmon Mute C4	903	840-898	
72	htrumpetax3	Trumpet A#3	1663	1602-1658	
73	htrumpetc3	Trumpet C3	1653	1598-1648	
74	htrumpetd2	Trumpet D2	1674	1628-1669	
75	htrumpetf3	Trumpet F3	1497	1454-1492	
76	htrumpetg2	Trumpet G2	1636	1598-1631	
77	jazzguitloop	Jazz Guitar Jazz Guitar Loop	119	9-114	

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
78	kotod3	Koto D3 Koto	5709	5666-5704	Only sample of that instrument
79	kpianob1	Piano B1	17232	10655-17227	
80	kpianob4	Piano B4	25217	15428-25212	
81	kpianob5	Piano B5	5399	2963-5394	
82	kpianocx4	Piano C#4	21574	21506-21569	
83	kpianodx5	Piano D#5	6146	6093-6141	
84	kpianof5	Piano F5	6980	6848-6974	
85	kpianof5 #02	Piano F5#02	6980	4148-4202	
86	kpianog2	Piano G2	22131	17637-22127	
87	lefone	Telephone	1585	8-1577	Based on name in 8MB E-mu bank
88	lowtumba	Low Tumba Tone	4022	7-4018	Based on name in 8MB E-mu bank
89	maracas	Maracas	3254	7-3250	
90	marimbac3	Marimba C3 Marimba	817	788-812	Only sample of that instrument
91	mbongotone	M Bongo Tone	2724	7-2720	Based on name in 8MB E-mu bank
92	mgtr	Muted Guitar Gtr Mute	836	766-831	Based on name in 8MB E-mu bank
93	nguitb2	Nylon Guitar B2 N Guitar B2	5193	5125-5188	Based on name in 8MB E-mu bank
94	nguitf2	Nylon Guitar F2 N Guitar F2	3829	3727-3824	Based on name in 8MB E-mu bank
95	oboeax3	Oboe A#3	998	958-992	
96	oboecx3	Oboe C#3	892	832-886	
97	oboefx3	Oboe F#3	1226	1177-1220	
98	Oboeresynwaved4	Oboe Resynth D4 Oboe Resynth	140	69-135	

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
99	ocarinafx2	Ocarina F#2 Ocarina	1187	1144-1182	Only sample of that instrument
100	octavewave	Octave Wave	140	69-135	
101	oohvoicec3	Ooh Voice C3 Ooh Voice	9102	35-9097	Only sample of that instrument
102	orchhit2	Orchestra Hit Orch Hit	4566	7-4562	Based on name in 8MB E-mu bank
103	organ19d4wave	Space Voice	140	69-135	The alias is because of its use in the 1MB E-mu bank.
104	organwave	Organ Wave Organ Wave 1	2675	2614-2669	
105	organwavea3	Organ Wave A3 Organ Wave 2	2940	2900-2934	
106	paisteping	Ride Cymbal Ride Ping	13293	7459-13288	Based on name in 8MB E-mu bank
107	pizzviolinc3	Pizzicato Strings C3 Pizzicato Strings Pizzicato Violin C3 Pizzicato Violin Pizz Violin C3 Pizz Violin	1560	1499-1555	Based on name in 8MB E-mu bank Only sample of that instrument
108	pluckharp	Pluck Harp Harp	3478	3409-3473	
109	quicadown	Quica Downstroke	1627	7-1623	Based on name in 8MB E-mu bank
110	quintoslap	Quinto Slap QuintoClosedSlap	2923	7-2919	Based on name in 8MB E-mu bank
111	quintotone	Quinto Tone	3091	7-3087	Based on name in 8MB E-mu bank
112	recorderax2	Recorder A#2 Recorder	1360	1298-1352	Only sample of that instrument
113	resynth4d4wave	Bowed Glass D4 Bowed Glass	140	69-135	Only sample of that instrument
114	rhodeschime	E Piano 1 Chime	284	7-279	
115	rideping	Ride Bell	6034	3614-6029	Based on name in 8MB E-mu bank
116	Sambawhistle	Samba Whistle	1687	1121-1673	

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
117	sawstackwavems	Saw Stack Wave Saw Stack	13749	289- 13745	
118	sb2	Slap Bass 2	2464	2345- 2459	
119	scratch	Scratch	1661	7-1657	
120	shakua2	Shakuhachi A2 Shakuhachi Siku E2	7468	2085- 7463	Only sample of that instrument
121	sikue2	Siku Bottle Blow E2 Bottle Blow	5314	7-5310	Only sample of that instrument
122	sinetick	Sinetick	73	7-68	
123	sinewave	Sine Wave Sample	140	69-135	Real time synthesis activated without "Sample" suffix
124	sitardx4	Sitar D#4 Sitar	2316	2275- 2311	Only sample of that instrument
125	slapbass1c3	Slap Bass 1 C3 Slap Bass 1	2121	1817- 2115	Only sample of that instrument
126	snare24	Snare	3877	7-3873	
127	squarewave	Square Wave Sample	2427	2365- 2420	Real time synthesis activated without "Sample" suffix
128	ssaxdx4	Soprano Sax D4 Soprano Sax	1189	1136- 1184	Only sample of that instrument
129	steeldrum	Steel Drum SteelDrum	2898	2857- 2891	Based on name in 8MB E-mu bank
130	stix	Drum Stick Side Stick	370	7-366	Based on name in 8MB E-mu bank
131	stringsdx4	Strings D#4	10727	3098- 10722	
132	stringsf3	Strings F3	8647	3415- 8642	
133	stringsg2	Strings G2	9309	2609- 9304	
134	synthbassloop	Synth Bass Loop Synth Bass	441	8-435	
135	Synthstringsc4	Synth Strings C4 Synth Strings	9967	1272- 9962	Only sample of that instrument

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
136	tamborine	Tambourine Brass Tambourine Tubular Bell D4	3604	2157-3585	Based on name in 8MB E-mu bank
137	tbelld4wave	Tubular Bell Tubular Bell Wave D4 Tubular Bell Wave	5417	4829-5412	Only sample of that instrument
138	timpani	Timpani Timp Drum	7699	7079-7695	Based on name in 8MB E-mu bank
139	triangle	Triangle Triangle Wave	2069	336-2064	Based on name in 8MB E-mu bank
140	troma3	Trombone A3	1334	1275-1329	
141	tromb2	Trombone B2	1331	1232-1326	
142	tromd4	Trombone D4	1121	1074-1116	
143	tromg4	Trombone G4	1569	1504-1564	
144	tubaax1	Tuba A#1 Tuba	1964	1831-1960	Only sample of that instrument
145	verbclickwave	Reverb Click Wave	1208	7-1204	
146	vibese2	Vibes E2 Vibes	782	696-777	Only sample of that instrument
147	vibraloop	Vibra Loop	788	9-783	
148	voxdodo	Doo Voice Oohs	3059	2999-3054	Based on name in 8MB E-mu bank
149	whitenoisewave	White Noise Sample	8294	7-8289	Real time synthesis activated without "Sample" suffix
150	woodblock	Wood Block Xylophone E4 Xylophone Xyo Unlooped E4	1164	7-1160	
151	Xyloe4looped	Xylo Unlooped XylophoneUnlooped E4 Xylophone Unlooped	980	7-976	Only sample of that instrument

Number	Original Sample Name	Acceptable Aliases	Sample length	Loop points	Description
152	xyloE4unlooped	Xylo Looped E4 Xylo Looped Xylophone Looped E4 Xylophone Looped	980	7-976	Only sample of that instrument

Sample specification is fixed at 44.1khz Mono with no links and tuning at 60 with no fine-tuning. You can either use discrete samples, or create your own SiliconSFe ROM containing the emulator samples.

Section 10

Compatibility information

10.1 Specification and structural compatibility

10.1.1 Legacy SF2.04 specification compatibility

SFe banks are not SoundFonts, and will not play properly on a legacy SF player! However, SFe banks are very similar to SoundFonts in terms of file structure, and should thus ideally be able to at least load on legacy SF players.

If a legacy SF player is not fully compatible with legacy SF2.04, then SFe files might not load at all.

10.1.2 INFO chunk and legacy compatibility

SFe files that use 32-bit chunk headers should be as compliant with the legacy standard as possible. SFSPEC24.PDF tells implementations to ignore additional sub-chunks in the INFO-list chunk, therefore it is possible to use the INFO-list chunk for additional data related to SFe.

If SFe files that use 32-bit chunk headers cannot be loaded on legacy soundcards, then something has went wrong.

SFe 4 does not support increased maximum length changes due to backward compatibility concerns.

10.1.3 phdr sub-chunk

This sub-chunk must contain at least two entries. Failure to do so will affect the compatibility with legacy SF players.

wBank changes are not recognised by legacy SF players. On such legacy players, only the m.s.b. is retained.

Presets with l.s.b.=0 should be the first presets internally. You may also want to use VArranger's workaround: for example, `115@ConcertGrand`.

10.1.4 Generators and modulators

In case of compatibility issues, the shAmount and wAmount options have been kept in for SFe 4. They may be removed in future versions of SFe.

If the SF version is below 2.1024, and the isng value is equal to EMU8000 or another E-mu sound engine:

- A 12dB low pass filter is used to ensure compatibility with the original AWE32 sound processor.
- Controller sources remain amplitude based.
- Ignore the whole modulator structure if a reserved source type is found.
- The Controller Source Type #16-#20 acts identically to the Controller Source Type #0-#4.

If the SF version is 2.04 or below, ignore the whole modulator structure if a reserved source type is found.

While default modulators 1-4 are not used in SFe version 4, SFe programs must still use them for older versions:

- If the SF version is 2.04, use the SF2.04 version of the Default Modulator 2.
- If the SF version is 2.01, use the SF2.01 version of the Default Modulator 2.
- If the sound engine is EMU8000, trigger legacy sound card mode.
- If the sound engine is E-mu 10K1, E-mu 10K2, X-Fi, or SFe 4, do not trigger legacy sound card mode.

Parameter units remain the same as SF2.04.

10.1.5 Error handling

Legacy SF players may halt on undefined chunks. Section 10.2 of SFSPEC21.PDF and SFSPEC24.PDF forbid the addition of sub-chunks to the sdta-list chunk, but Creative/E-mu themselves decided to do it anyway when developing SF2.04 (by using the sm24 sub-chunk).

Legacy SF players may halt on unknown enums, which goes against the legacy SF2.04 specification. However, at least some legacy soundcards do not error out on an unknown enum.

10.2 Sample compatibility

10.2.1 32-bit samples

Players fully compliant with legacy SF2.04 should be able to play files with 32-bit samples at 24-bit quality. However, these files may fail on software (such as Polyphone 2.4.x) that looks specifically for `sm24` and reject anything else. It is therefore not recommended to use 32-bit samples with legacy SF2.0x players.

The `sm32` sub-chunk is implemented in the same way as `sm24` was by E-mu, to maximise compatibility, but due to the massive size, un practicality and compatibility implications of 32-bit samples, we recommend that you use the `sm32` sub-chunk only with 64-bit chunk headers.

10.2.2 8-bit samples

The legacy specification instructs programs to ignore an `sm24` sub-chunk that is not paired with an `smp1` sub-chunk. Software that is compatible with legacy SF2.04, but is incompatible with SFe, will not play 8-bit samples, as legacy software cannot read the `ISFe-info` sub-chunk. Legacy sound cards will also ignore lone `sm24` chunks.

10.2.3 ROM samples

ROM samples may still be used; therefore bit 16 should not be ignored.

10.2.4 Proprietary compression formats

If an incompatible proprietary compression format is found, then the program may offer decompression of the incompatible file before use if you have permission from the original author of the compression formats. SF compression programs are notorious for restrictive licensing.

Section 11

Program specification

11.1 Program compatibility levels

These are guidelines to help SFe program developers write their programs. Not all of them need to be followed, but doing so will ensure that your program works with all SFe 4.0 banks.

All SFe programs must meet level 1 requirements. If they do not, guarantees that an SFe file will work on the program can not be made.

Level 1 is the level that embedded applications should be expected to include. Level 2 is the minimum 32-bit level for computers. Level 3 is the recommended 64-bit level for a good experience. Level 4 is maxed out. Requirements for levels 1-4 will increase with new versions.

If an implementation is unable to reach the layering requirements without crashing, it does not entirely meet the level.

11.1.1 File format specifications

	Level 1	Level 2	Level 3	Level 4
File size representation	Unsigned 32-bit integer You must not use signed integers	Unsigned 32-bit integer You must not use signed integers	Unsigned 64-bit integer You must not use signed integers	Unsigned 64-bit integer You must not use signed integers
File size limit	System memory	Based on chunk header type	Based on chunk header type	Based on chunk header type
Sample streaming	System memory	System memory and disk streaming	System memory and disk streaming	System memory and disk streaming
Total file size limit	System memory	At least 32 GiB	No limit	No limit
Multiple files	Optional	8 or more	256 or more	No limit
Legacy support	Full quality: SF2.01 and Werner SF3 Playback: SF2.04	Full quality: SF2.01 and Werner SF3 Playback: SF2.04	Full quality: SF2.01, SF2.04 and Werner SF3	Full quality: SF2.01, SF2.04 and Werner SF3
Header support	32-bit static Werner SF3 format Uncompressed, OGG Proprietary formats forbidden	32-bit static Werner SF3 format Uncompressed, OGG Proprietary formats forbidden	32-bit static, 64-bit static Werner SF3 format Uncompressed, OGG Proprietary formats forbidden	32-bit static, 64-bit static, dynamic, RIFX Werner SF3 format Uncompressed, OGG Proprietary formats forbidden
Sample compression	SFe: .sft SF2.0x: .sf2 Werner SF3: .sf3 Any other uncompressed format is allowed	SFe: .sft SF2.0x: .sf2 Werner SF3: .sf3 Any other uncompressed format is allowed	SFe: .sft SF2.0x: .sf2 Werner SF3: .sf3 Any other uncompressed format is allowed	SFe: .sft SF2.0x: .sf2 Werner SF3: .sf3 Any other uncompressed format is allowed
File extension	New chunks, feature flags	New chunks, feature flags	New chunks, feature flags	New chunks, feature flags
Information/ Metadata				

11.1.2 Sample specifications

	Level 1	Level 2	Level 3	Level 4
Maximum sample rate	44 100 Hz or greater	50 000 Hz or greater	96 000 Hz or greater	No limit
Sample bit depth	16-bit or greater Ignore unsupported sdta sub-chunks	16-bit or greater Ignore unsupported sdta sub-chunks	24-bit or greater Ignore unsupported sdta sub-chunks	32-bit or greater Ignore unsupported sdta sub-chunks
Individual sample length	24-bit or greater	32-bit or greater	32-bit or greater	Based on chunk header type
Loop point sets	1	1	1	1
Sample linking	Mono, Left/Right Includes ROM samples and SFe Compression	Mono, Left/Right, "Link" Includes ROM samples and SFe Compression	Mono, Left/Right, "Link" Includes ROM samples and SFe Compression	Mono, Left/Right, "Link" Includes ROM samples and SFe Compression
Number of channels	Mono, Stereo	Mono, Stereo	Mono, Stereo	Mono, Stereo
Sample name length	Display 8 characters Write 20 characters	Display 20 characters Write 20 characters	Display 20 characters Write 20 characters	Display 20 characters Write 20 characters

11.1.3 Instrument specifications

	Level 1	Level 2	Level 3	Level 4
Samples per instrument	64 or greater	4096 or greater	16384 or greater	No limit
Simultaneous sample playback	2 or greater	32 or greater	64 or greater	No limit
Modulators	All 4.0 modulators	All 4.0 modulators	All 4.0 modulators	All 4.0 modulators
Default modulators	All 4.0 default modulators	All 4.0 default modulators	All 4.0 default modulators	All 4.0 default modulators
Enumerators	All 4.0 enumerators	All 4.0 enumerators	All 4.0 enumerators	All 4.0 enumerators
Instrument name length	Display 8 characters Write 20 characters	Display 20 characters Write 20 characters	Display 20 characters Write 20 characters	Display 20 characters Write 20 characters

11.1.4 Preset specifications

	Level 1	Level 2	Level 3	Level 4
Instruments per preset	2 or greater	16 or greater	64 or greater	No limit
Simultaneous sample playback	4 or greater	256 or greater Or up to polyphony limit	512 or greater Or up to polyphony limit	No limit
Preset name length	Display 8 characters Write 20 characters	Display 20 characters Write 20 characters	Display 20 characters Write 20 characters	Display 20 characters Write 20 characters

11.1.5 Player specifications

	Level 1	Level 2	Level 3	Level 4
Polyphony limit	32 or greater with mono samples 16 or greater with stereo samples	256 or greater with mono samples 128 or greater with stereo samples	512 or greater with mono samples 256 or greater with stereo samples	No limit
Percussion toggle	byBankMSB/ byBankLSB bit 1	byBankMSB/ byBankLSB bit 1	byBankMSB/ byBankLSB bit 1	byBankMSB/ byBankLSB bit 1
Control change 0/32 (Bank select)	CC0: byBankMSB bits 2-8 CC32: byBankLSB bits 2-8	CC0: byBankMSB bits 2-8 CC32: byBankLSB bits 2-8	CC0: byBankMSB bits 2-8 CC32: byBankLSB bits 2-8	CC0: byBankMSB bits 2-8 CC32: byBankLSB bits 2-8
Control change 7/39/8				
Control change 40/10/42	Mandatory	Mandatory	Mandatory	Mandatory
Control change 1/33/2/34/4/36	CC1/33 (modulation wheel) required	CC1/33 (modulation wheel) required	Mandatory	Mandatory
Control change 5/37/11/43/84				
Control change 6/38	GM1 data entry required	GM1 data entry required	GM1 data entry required	GM1 data entry required
Control change 12/13/44/45				
Control change 16/17/18/19				
Control change 48/49/50/51				
Control change 80/81/82/83				
Control change 64/66/67	Mandatory	Mandatory	Mandatory	Mandatory
Control change 68/69	Optional	Optional	Mandatory	Mandatory

	Level 1	Level 2	Level 3	Level 4
Control change	CC70 variation CC71 timbre	CC70 variation CC71 timbre	CC70 variation CC71 timbre	CC70 variation CC71 timbre
70/71/72/73/74	CC72 release time CC73 attack time			
Control change	CC74 brightness CC75 decay time CC76 vibrato rate			
75/76/77/78/79	CC77 vibrato depth CC78 vibrato delay			
Control change 91/93	Reverb and chorus	Reverb and chorus	Reverb and chorus	Reverb and chorus
Control change 92/94/95	GM1 level effects unit required			
Control change 96/97/98	Mandatory	Mandatory	Mandatory	Mandatory
Control change 99/100/101				
Control change 120/121/122/1				
23	Mandatory	Mandatory	Mandatory	Mandatory
Control change 124/125/126/1				
27				
GM1 reset	Mandatory	Mandatory	Mandatory	Mandatory
SMF formats	0 and 1	0 and 1	0, 1 and 2	Any
Number of channels	16 or greater	16 or greater	64 or greater	No limit
ROM emulator	ROM sample support or emulator optional	ROM sample support or emulator optional	ROM sample support or emulator required	ROM sample support or emulator required

11.2 Converting between legacy SF and SFe

11.2.1 Conversion from legacy SF2.04 to SFe

- Upgrade the `ifil` version in the header from `wMajor=2, wMinor=4` to `wMajor=2, wMinor=1024`.
- Overwrite the `isng` value with SFe 4.
- Create an ISFe-list sub-chunk with information: `SFty = "SFe-static", SFvx = 4, 0, Release Candidate, 2, "4.0-rc2"`, flag corresponding to features used in the bank.

11.2.2 Conversion from SFe to legacy SF2.04

- Downgrade the ``ifil` version in the header from `wMajor=2, wMinor=1024` to `wMajor=2, wMinor=4`.
- Overwrite the `isng` value with X-Fi.
- Decompress the samples if the bank uses SFe Compression.
- Downsample any 32-bit samples to 24-bit, and remove the `sm32` sub-chunk.
- Convert all 8-bit samples to 16-bit by moving the sample data to the `smpl` sub-chunk.
- If the first 8 bits of either `wPreset` or `wBank` are identical to another patch, keep only the patch with bits 9-16 clear. If this is not found, keep only the last of the patches in the record. Clear bits 9-16 of all patches afterwards.
- Remove the ISFe-list sub-chunk.

11.2.3 Conversion from SFe to legacy SF2.01

- Downgrade the `ifil` version in the header from `wMajor=2, wMinor=1024` to `wMajor=2, wMinor=1`.
- Overwrite the `isng` value with `EMU8000` or `E-mu 10K1`.
- Decompress the samples if the bank uses SFe Compression.
- Downsample any 24-bit or 32-bit samples to 16-bit, and remove the `sm24` and `sm32` sub-chunks.
- Convert all 8-bit samples to 16-bit by moving the sample data to the `smp1` sub-chunk.
- If the first 8 bits of either `wPreset` or `wBank` are identical to another patch, keep only the patch with bits 9-16 clear. If this is not found, keep only the last of the patches in the record. Clear bits 9-16 of all patches afterwards.
- Remove the `ISFe-list` sub-chunk.

11.3 Converting between chunk header types

11.3.1 Conversion of 32-bit static to 64-bit static

- Replace RIFF with RF64.
- Add a `ds64` chunk with chunk size.
- Set `ckSize` to `FF FF FF FF`.
- Replace `sfbk` with `sfen`.

11.3.2 Conversion of 64-bit static to 32-bit static

- Make sure file size does not exceed 4GiB.
- Set `ckSize` to the value of `ds64`.
- Delete `ds64`.
- Replace RF64 with RIFF.
- Rename `sfen` to `sfbk`.

11.3.3 Conversion between 32-bit static and RIFF

- Replace RIFF with RIFX (or vice versa).
- Swap the bytes to convert between big endian and little endian.

11.4 File repair programs

11.4.1 Repairing Structurally Unsound errors

ifil sub-chunk errors

To fix ifil sub-chunk errors, file repair programs must determine the correct version number of the SFe program by inspecting the file structure. File repair programs should never simply add the current SFe specification version, as that can cause a non-critical error relating to an ifil version mismatch due to the specification version being written in ISFe-list (SFvx) and not ifil.

If the file is too damaged to determine the correct ifil version, then the repair program should repair these problems first. If, after this, a definitive SF version cannot be determined, the user should be given the option to manually enter a new SF file version. The correct data in ISFe-list should also be included if the file is an SFe bank.

Alternatively, file repair programs can allow the user to manually enter the SF file version without allowing the program to automatically determine the necessary version.

PHDR sub-chunk errors

In the case of a missing PHDR sub-chunk, the program should just reconstruct an "empty" PHDR sub-chunk. The bank will not be usable until the PHDR sub-chunk is manually re-created.

If the PHDR sub-chunk contains at least two records, but is not a multiple of 38 bytes, then this indicates a structural error in one or more records. Each record can be inspected and the structure repaired. The user should also be given the choice between repairing the records and just deleting them. Repaired records must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

Non-monotonic preset bag indices should be reordered if encountered.

The PBAG sub-chunk size is usually corrected by correcting the incorrect value(s). The correct value can be reconstructed once the preset bag indices have been reordered.

PBAG sub-chunk errors

In the case of a missing PBAG sub-chunk, the program should just reconstruct an "empty" PBAG sub-chunk. The bank will not be usable until the PBAG sub-chunk is manually re-created.

If the PHDR sub-chunk is not a multiple of 4 bytes, or if there is a mismatch between the generator or modulator indices and corresponding PGEN/PMOD sub-chunk sizes, then this indicates a structural error in one or more of the preset zones. Each zone can be inspected and the structure repaired. The user should also be given the choice between repairing the zones and just deleting them. Repaired zones must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

Non-monotonic generator or modulator indices should be reordered if encountered.

The PGEN/PMOD sub-chunk sizes are usually corrected by correcting the incorrect value(s). The correct value can be reconstructed once the preset bag indices have been reordered.

PMOD sub-chunk errors

In the case of a missing PMOD sub-chunk, the program should just reconstruct an "empty" PMOD sub-chunk. The bank may be usable, but there will be a significant quality loss until the PMOD sub-chunk is manually re-created.

If the PMOD sub-chunk is not a multiple of 10 bytes, then this indicates a structural error in one or more of the preset modulators. Each modulator can be inspected and the structure repaired. The user should also be given the choice between repairing the modulators and just deleting them. Repaired modulators must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

PGEN sub-chunk errors

In the case of a missing PGEN sub-chunk, the program should just reconstruct an "empty" PGEN sub-chunk. The bank will not be usable until the PGEN sub-chunk is manually re-created.

If the PGEN sub-chunk is not a multiple of 4 bytes, then this indicates a structural error in one or more of the preset generators. Each generator can be inspected and the structure repaired. The user should also be given the choice between repairing the generators and just deleting them. Repaired generators must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

The instrument generator value is usually corrected by correcting the incorrect value(s) by inspecting the structure and determining the correct value for the instrument generator value or terminal instrument.

INST sub-chunk errors

In the case of a missing INST sub-chunk, the program should just reconstruct an "empty" INST sub-chunk. The bank will not be usable until the INST sub-chunk is manually re-created.

If the INST sub-chunk contains at least two records, but is not a multiple of 22 bytes, then this indicates a structural error in one or more records. Each record can be inspected and the structure repaired. The user should also be given the choice between repairing the records and just deleting them. Repaired records must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

Non-monotonic instrument bag indices should be reordered if encountered.

The INST sub-chunk size is usually corrected by correcting the incorrect value(s). The correct value can be reconstructed once the instrument bag indices have been reordered.

IBAG sub-chunk errors

In the case of a missing IBAG sub-chunk, the program should just reconstruct an "empty" IBAG sub-chunk. The bank will not be usable until the IBAG sub-chunk is manually re-created.

If the IHDR sub-chunk is not a multiple of 4 bytes, or if there is a mismatch between the generator or modulator indices and corresponding IGEN/IMOD sub-chunk sizes, then this indicates a structural error in one or more of the instrument zones. Each zone can be inspected and the structure repaired. The user should also be given the choice between repairing the zones and just deleting them. Repaired zones must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

Non-monotonic generator or modulator indices should be reordered if encountered.

The IGEN/IMOD sub-chunk sizes are usually corrected by correcting the incorrect value(s). The correct value can be reconstructed once the instrument bag indices have been reordered.

IMOD sub-chunk errors

In the case of a missing IMOD sub-chunk, the program should just reconstruct an "empty" IMOD sub-chunk. The bank may be usable, but there will be a significant quality loss until the IMOD sub-chunk is manually re-created.

If the IMOD sub-chunk is not a multiple of 10 bytes, then this indicates a structural error in one or more of the instrument modulators. Each modulator can be inspected and the structure repaired. The user should also be given the choice between repairing the modulators and just deleting them. Repaired modulators must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

IGEN sub-chunk errors

In the case of a missing IGEN sub-chunk, the program should just reconstruct an "empty" IGEN sub-chunk. The bank will not be usable until the IGEN sub-chunk is manually re-created.

If the IGEN sub-chunk is not a multiple of 4 bytes, then this indicates a structural error in one or more of the instrument generators. Each generator can be inspected and the structure repaired. The user should also be given the choice between repairing the generators and just deleting them. Repaired generators must be indicated so the user can make any required adjustments. This does not have to be in the file, rather just in the user interface of the program.

The sampleID generator value is usually corrected by correcting the incorrect value(s) by inspecting the structure and determining the correct value for the sampleID generator value or terminal sampleID.

SHDR sub-chunk errors

If there is no usable sample data, the program should just reconstruct an "empty" SHDR sub-chunk. However, if there is usable sample data, the program may attempt to partially reconstruct the SHDR sub-chunk. User input will be required.

If there is no valid irom sub-chunk, the program can clean these samples if they aren't being used in the bank. If they are used by the bank, then the program should then attempt to read the sample as a non-ROM sample, and if the samples are valid, then this means that the problem is fixed. If there is no valid sample, then either a new sample must be provided by the user, or a ROM should be defined for the irom sub-chunk.

Unknown sub-chunk errors

Unknown sub-chunk errors can easily be fixed; the unknown sub-chunks are simply removed.

Compressed sample errors

In SFe Compression, all PCM samples must go before compressed samples. This is fixed by rearranging the sample data in the smp1 chunk, and then updating the shdr sub-chunk.

ckSize errors

ckSize errors are easily solved by checking the ckSize values and fixing them with the correct file sizes, and in the case of 64-bit static shunk headers, setting the ds64 chunk size to the correct value.

Incompatible chunk header errors

Incompatible chunk header errors are solved by replacing chunk headers with a compatible version, or swapping the endianness of the data.

Incompatible SFty errors

TSC-related errors are corrected by rearranging the chunks, while 8-bit samples can be converted to 16-bit.

8-bit sample errors

8-bit sample errors are corrected by attempting to repair the `smpl` sub-chunk. If there is no `smpl` sub-chunk, then an SFe player can attempt to load the file with 8-bit sample mode.

11.4.2 Repairing non-critical errors

isng sub-chunk errors

In the case of a missing `isng` sub-chunk, a value of SFe 4 should be written.

If it doesn't end in a zero-valued byte, then add a zero-valued byte and warn the user to verify if the value is correct.

ICRD sub-chunk errors

In the case of a missing `ICRD` sub-chunk, a value corresponding to the current ISO-8601 date and time should be written.

Should the `ICRD` sub-chunk not be a valid ISO-8601 date or time, firstly attempt to parse the current value. For example, it may be in a different language. If it can't be parsed, then write a value corresponding to the current ISO-8601 date and time.

INAM, IENG, IPRD, ICOP, ICMT or ISFT sub-chunk errors

Missing sub-chunks should be filled in intuitively. For example, for an INAM sub-chunk, the file name or Unnamed bank are good starting points.

If it doesn't end in a zero-valued byte, then add a zero-valued byte and warn the user to verify if the value is correct.

With the exception of the INAM chunk, if the sub-chunk's value cannot be faithfully copied as a string, the user should be given a chance to add the correct value in UTF-8.

irom or iver sub-chunk errors

Missing sub-chunks should not be filled in unless there are ROM samples that are linked.

smpl, sm24 and sm32 sub-chunk errors

If both the sm24 and sm32 sub-chunks are present, then these sub-chunks should be combined into 16-bit samples in one smpsub-chunk. Because endianness may not be clear in this situation, you must allow the user to select the endianness of the operation.

If there is a mismatch in the ifil version and the presence of both sm24 and sm32 sub-chunks, then allow the user to select whether they want to keep these sub-chunks. If there is only one such sub-chunk, then allow the user to update the bank to use SFe with 8-bit samples.

PHDR sub-chunk errors

Any dwLibrary, dwGenre or dwMorphology value should be cleared because current versions of SFe do not implement this yet. Presets without any zones should be given a zone named Empty preset or similar.

PBAG sub-chunk errors

These problems can easily be fixed by giving zones after the first (including the global zone) an instrument generator.

PMOD or IMOD sub-chunk errors

To fix unknown or inappropriate "link" values, they should be corrected to a value given by the user. Out-of-range modulator indexes should be fixed by letting the user select the correct modulator, or define a new modulator if it doesn't exist. Circularly linked modulators should be highlighted to allow the user to break the circular links. Modulators with the same enumerators should also be highlighted, so the enumerators can be changed by the user.

PGEN sub-chunk errors

Unknown sfGenOper values should be corrected to a value given by the user. Modulators with the same sfGenOper enumerator as another modulator should be highlighted, so it can be changed by the user. Generators and modulators in inappropriate places should also be moved with references corrected. Non-global lists that don't have an instrument generator at the end should have an instrument generator added.

INST sub-chunk errors

To fix this issue, the instrument should be deleted unless it's being used by a preset. If so, the instrument should be kept, but the user should be warned so they can fix the problem.

IBAG sub-chunk errors

These problems can easily be fixed by giving zones after the first (including the global zone) a sampleID generator.

IGEN sub-chunk errors

Unknown sfGenOper values should be corrected to a value given by the user. Modulators with the same sfGenOper enumerator as another modulator should be highlighted, so it can be changed by the user. Generators and modulators in inappropriate places should also be moved with references corrected. Non-global lists that don't have a sampleID generator at the end should have an sampleID generator added.

SHDR sub-chunk errors

If the sample rate is zero, then it should be corrected automatically to the correct value via automatic pitch detection and then highlighted for the user to verify. Bad original pitch values should be corrected to the value with automatic pitch detection. Non-zero `wSampleLink` values should be set to zero if SFe Compression is in use.

You should not edit the sample rate unless it is zero.

Undefined and unknown enum, palette value and source type errors

Anything that uses these values should either be removed or highlighted so the user can inspect it. If a value is undefined for the SFvx version, but defined for a newer version, then the program can also offer to update the version to a newer version.

Precedence errors

In this case, the user should be allowed to change the generator value or to remove the offending generator entirely.

Parameter value and padding errors

In this case, the illegal or missing parameter value should be shown to the user, so they can rectify it. Padding should automatically be added wherever possible.

Unknown chunk errors

If a chunk is undefined in the SFvx version given, but defined in a later version, then the SFvx and/or if1 versions should be updated. If it is undefined in any version, then it can safely be deleted.

Compression errors

The user should be given the choice of compression algorithms to use for the sample, and `wSampleLink` values should be set to zero.

ISFe-list chunk errors

The ISFe-list sub-chunk can be repaired, however information in the ISFe-list sub-chunk should be reconstructed.

ifil chunk errors

If features or feature flags from a newer ifil or SFvx version of SFe are found on a bank with an older declared version of SFvx, then the version should be updated to the detected version of SFvx.

Proprietary compression errors

If a proprietary compression format is supported by the SFe program, then it should be decompressed and automatically recompressed into a lossless format using SFe Compression.

wPreset value errors

If invalid wPreset values are found, then they should be highlighted by the program, and the user should have the opportunity to select a different wPreset value.

Duplicated preset location errors

Duplicated preset locations should be highlighted, and the user should have the opportunity to select a different preset value.

File size limit errors

32-bit static headers and structures can be replaced with 64-bit counterparts if the file size was found to exceed 4 GiB. Alternatively, if TSC mode is supported by the SFe program, then it can be activated by moving the sdt-a-info chunk to the end and setting the correct feature flag in the ISFe-list sub-chunk.

Feature flag errors

Feature flags can be updated to accurately reflect the features that the bank uses.

11.4.3 Manual repair

Manual repair is used to completely fix a broken SFe bank. It can be partially automated, however some parts of manual repair are not automatable, and the parts that aren't automatable are highlighted.

Manual repair is intended to be part of SFe bank development tools such as SFe editors or dedicated file repair programs. It is also suitable for situations where control over the repair process by the user is required. However, it is not suitable for SF players.

Manual repair can be used to repair all Structurally Unsound errors and non-critical errors listed in section 11.4. This is the main advantage that it holds over automatic repair.

You can use it to fix a corrupted bank. A manual repair program is very useful if you develop banks, as your data may be damaged at any time for any reason.

11.4.4 Automatic repair

Automatic repair allows SFe players to play some "Structurally Unsound" banks by automatically repairing them at load time. It consists of a partial file repair program built directly into the SFe player. With future versions of this specification, SFe players that only support a higher version of the format will also gain the ability to translate old banks to run properly.

Automatic repair is intended to be part of SFe players. It is an aid to help SFe players play banks that are otherwise "Structurally Unsound". It is not a substitute for repairing the bank using file repair programs. File repair programs should always include an option to use manual repair.

Automatic repair fixes structural defects in SFe and legacy SF2.04 banks seamlessly and transparently. It should automatically repair all Structurally Unsound errors listed in section 11.4.1 except for compressed sample errors, 8-bit sample errors and anything that requires user input.

It should also repair all non-critical errors listed in section 11.4.2 except for compression errors, wPreset value errors, file size limit errors and anything that requires user input. The repair of proprietary compression errors is optional and contingent on support for such formats.

The program developers are allowed to use any repair strategy listed in section 11.4.

Automatic repair requires that the implementation give a message when invoked. This is to encourage bank developers not to rely on the feature to implicitly define any parameters.

Automatic repair is a great method to fix issues that prevent SF banks from loading correctly, but it can't repair all defects. For example, anything that requires user input is outside of the scope of automatic repair, and should be dealt with by a file repair program or SFe editor program instead.

This feature must also never overwrite files. This is the job of a file repair program. If system memory is low, it is acceptable for automatic repair implementations to create a patched bank in temporary file storage.

You must correctly setup your banks to ensure that they run properly; do not use automatic repair to implicitly define certain parameters. If you do so, then your bank may not run properly on a legacy SF player or an SFe player that doesn't implement automatic repair.

11.5 Why these guidelines?

11.5.1 File Size Representation

- SFe is a 32-bit or 64-bit format.
- Incorrect numbers of bits will therefore result in program non-compliance.
- Signed integers are not useful because a negative file size is impossible.
- To ensure that file size limits are not arbitrarily "cut in half", signed integers are prohibited.
- Programs using signed integers to represent file size are therefore not compliant with SFe.

11.5.2 File Size Limit

- You should not impose additional file size limits in SFe programs.
- SFe files that exceed these limits may not play.
- Because compatibility is not guaranteed, such programs may not be SFe compliant.

11.5.3 Sample Streaming

- You should strive to include disk streaming support.
- This is because the file size limit is almost always greater than the system memory installed.
- System memory streaming options may be available for higher performance.

11.5.4 Total File Size Limit and Multiple Files

- Always provide space for more than one file.
- SFe allows flexibility in implementing multiple files.
- It will be standardised in the future depending on the most popular system.
- The program specification requires multiple files, to guarantee proper operation of "split bank" files.
- Not implementing multiple file operation will affect compatibility.
- Such programs therefore may not be SFe compliant.

11.5.5 Legacy Support

- SFe is designed to reconcile incompatible SF extensions.
- Werner SF3 support is required; the SFe format incorporates Werner SF3 structures (SFe Compression).
- Programs that cannot recognise Werner SF3 format (SFe Compression) may not be able to open SFe files.
- 24-bit support is optional, but if 24-bit is unsupported in an SFe player, legacy SF2.04 files must not be rejected.
- Backward compatibility with legacy SF2.0x is sacrosanct when using 32-bit chunk headers, take care when creating a program.
- "Synthfont Custom Features" support is not currently included in the specification.

11.5.6 Header Support

- The header should be RIFF and sfblk with 32-bit chunk headers to preserve legacy SF2.0x compatibility.
- 32-bit only programs should recognise 64-bit chunk headers and give an error.
- The header should be RF64 and sfen with 64-bit chunk headers.
- All programs that support 64-bit chunk headeres must also support 32-bit chunk headers.

11.5.7 Sample Compression

- Compression algorithms vary, but we recommend at least one of each lossy and lossless algorithms.
- OGG and FLAC are specified in the specification.
- Werner SF3 was designed to use OGG files, so OGG is a good start for lossy algorithms.
- Opus is an alternative that is "better" than OGG according to some people.
- FLAC and "Wavpack" are free lossless compression algorithms for sounds.
- Note that while "Wavpack" has advantages over FLAC, it is not widely supported.

11.5.8 File Extension, Structure, Information and Metadata

- The file extension is .sft. Chunk headers must be detected.
- Do not save SFe files with an extension .sf2, as it may confuse legacy SF2.0x players. Remember that SFe files are not SoundFonts!
- Use the ifil and SFvx sub-chunks to determine version, do not use the file extension.
- For SFe, you must use the prescribed chunk sizes and limits in the specification.

11.5.9 Sample Specifications

- Level 3 requires 96kHz frequency, because it is the next standard frequency above 50kHz as specified by SF2.04.
- 88.2kHz is not recommended because it is non-standard.
- Stray sdt_a sub-chunks must be ignored. Erroring out on extra sdt_a chunks is not compliant with SF2.04.
- There must not be additional limitations to sample length, besides the file size limits.
- Sample linking features from legacy SF2.04 must be supported (with the exception of SFe Compression).

11.5.10 Instrument Specifications

- High numbers of samples per instrument are specified to ensure that the program works efficiently.
- Ideally, programs should support unlimited numbers of samples per instrument.
- You must not add restrictions to number of samples per instrument or simultaneous samples.
- Modulators are fixed and must be implemented as shown in the specification and not as in SF2.04.
- In legacy SF2.04 compatibility mode, bug compatibility can be implemented.
- When field sizes increase from 20 characters, programs should be able to display at least 64 characters.
- If it is not possible to display the entire field, use an ellipsis or similar.

11.5.11 Player Specifications

- A minimum polyphony of 256 notes is in place if possible.
- The byBankLSB percussion toggle is to fix the bank select LSB function.
- All control changes listed as "mandatory" must be supported.
- If it is not supported, playback may be severely impacted, resulting in a program that is not SFe compliant.
- This is already a problem with legacy SF2.0x, as not all SF2 players support modulators.
- Reset support and multiple simultaneous drum kits might not be required, but it is suggested.
- 64-channel MIDI file support is suggested, and may be required in the future for SFe.
- The ROM emulator may be required in the future, as many SFe files will make full use of it.

11.6 How to test your program with SFSpecTest

11.6.1 What does SFSpecTest do?

By using SFSpecTest by mrbumpy409 ([available here](#)), you can test your SFe player and determine which feature flags to set.

SFSpecTest was written by the author of GeneralUserGS, one of the most popular legacy SF2.04 banks, and is thus a good benchmark for legacy SF2.04 players.

Because SFe is a superset of legacy SF2.04, it is also a good tool to determine what SF2.04 features your program supports, allowing you to set the correct feature flags for your program.

11.6.2 Branch 00 Foundational synthesis engine

In leaf 00:00, if your player passes SFSpecTest test #8 (Scale Tune/Root Key), you can set all four defined bits.

In leaf 00:03, set bits 1-16 to the maximum frequency attained in SFSpecTest test #9 (Initial Filter Cutoff), and bits 17-24 to the maximum resonance attained in SFSpecTest test #10 (Filter Resonance).

In leaf 00:04, if your player passes SFSpecTest test #11 (Attenuation Amount), you can set the first two defined bits. If your player passes SFSpecTest test #5a (Modulation LFO A) and #12 (Negative Attenuation Amount), you can also set the third bit.

In leaf 00:05, set bit 1 if you pass SFSpecTest test #17a (Reverb A), bit 2 if you pass SFSpecTest test #17b (Reverb B), bit 3 if you pass SFSpecTest test #17c (Reverb C), bit 9 if you pass SFSpecTest test #18a (Chorus A), bit 10 if you pass SFSpecTest test #18b (Chorus B), and bit 11 if you pass SFSpecTest test #18c (Chorus C). Set bit 4 if your reverb can be adjusted, and set bit 12 if your chorus can be adjusted.

In leaf 00:06, if your player passes SFSpecTest test #5 (Modulation LFO), you can set bit 4. If your player passes SFSpecTest test #6 (Vibrato LFO) and test #7 (Mod Wheel to LFO), you can set bit 2.

In leaf 00:07, if your player passes SFSpecTest test #1 (Volume Envelope), you can set bits 1-6. If your player passes SFSpecTest test #2 (Modulation Envelope), you can set bits 9-14. If your player passes SFSpecTest test #3 (Key Number to Decay), you can set bits 8 and 16. If your player passes SFSpecTest test #4 (Key Number to Hold), you can set bits 7 and 15.

In leaf 00:0a, set bit 3 if you pass SFSpecTest test #21 (Exclusive Class). Set bit 6 if you pass SFSpecTest #16 (Sample Offset).

11.6.3 Branch 01 Modulators and NRPN

In leaf 01:01, set bit 5 if you pass SFSpecTest test #20a (Pitch Bend A) and test #20b (Pitch Bend B). Set bit 6 if you pass SFSpecTest test #20c (Pitch Bend C).

In leaf 01:02, set bit 1 if you pass SFSpecTest test #15 (CC1 to Filter Cutoff).

In leaf 01:06, set bit 1 if you pass SFSpecTest test #13 (Velocity to Attenuation Curve), and set bit 2 if you pass SFSpecTest test #14a (Velocity to Initial Filter Cutoff Curve A) and test #14b (Velocity to Initial Filter Cutoff Curve B). Set bit 4 if you pass SFSpecTest test #15 (CC1 to Filter Cutoff). Set bit 17 if you emulate SF2.00 behaviour as shown in test #14c (Velocity to Initial Filter Cutoff Curve C), set bit 18 if you emulate SF2.01 behaviour as shown in test #14d (Velocity to Initial Filter Cutoff Curve D), and set bit 19 if you emulate SF2.04 behaviour as seen in test #14e (Velocity to Initial Filter Cutoff Curve E).

11.7 Courtesy actions

For the benefit of the SFe community, please:

- share all modifications implemented in a program under this license
- do not remove the link to the latest version of the specification

The above requirements are not required due to requirements listed by the Open Source Definition, but are good practice when redistributing the specification.

Section 12

Glossary

This glossary is broadly the same as the glossary in SFSPEC24.PDF, with these additions and changes:

- Articulation – Modulation of available parameters and usage of extra samples to produce expressive musical notes.
- AWE64 – The successor to the famous AWE32, adding features such as waveguide synthesis. Used the EMU8000 synthesizer chip, like the preceding AWE32. Available in "Value" or "Gold" versions.
- Branch – A subdivision of a tree structure containing either sub-branches or leaves that include values.
- BW64 – Broadcast Wave 64, used in the RF64 Header.
- Case-insensitive – Indicates that a UTF-8 character or string treats alphabetic characters of upper or lower case as identical.
- Case-sensitive – Indicates that a UTF-8 character or string treats alphabetic characters of upper or lower case as distinct.
- Cognitone SF4 – An incompatible modification to Werner SF3 to allow support for FLAC audio compression. Because it is considered proprietary compression, usage is not allowed in SFe.
- DAHDSR – Stands for Delay, attack, hold, decay, sustain, release. The six-step envelope system used in SF and SFe.
- Downloadable – legacy SF2.0x, Werner SF3 or SFe file obtained from the internet. (Old meaning referred to the obsolete ROM system)
- EMU10K1 – The successor to the EMU8000, designed by E-mu® for the Creative Labs SB Live!.
- EMU10K2 – An update to the EMU10K1, designed by E-mu® for the Creative Labs SB Audigy.
- EMU20K1 – The successor to the EMU10K2, designed by E-mu® for the Creative Labs SB X-Fi.
- EMU20K2 – An update to the EMU20K1, please refer [here](#) for information on SB X-Fi cards that include it.
- FLAC – A lossless audio compression format commonly used in open-source software. Supported by Werner SF3, but not commonly used for that purpose.
- Hold – The portion of the DAHDSR envelope after the attack portion, but before the decay portion starts.
- Leaf – A value found in a tree structure at the end of a branch.
- Legacy sound card – A Sound Blaster® (or other sound card) that uses a hardware MIDI synthesiser capable of using banks in the SoundFont® format.
- Lossless compression – Said of a compression format that retains all of its data when compressed. In terms of audio, there is no loss in quality in losslessly compressed audio.

- Lossy compression – Said of a compression format that does not retain all of its data when compressed. In terms of audio, there is a loss in quality in lossily compressed audio.
- MIDI Bank – Groups of up to 128 presets, which can be selected by the two MIDI "Bank Select" control changes (CC00 and CC32).
- OGG – See "Vorbis".
- Opus – A lossy audio compression format, slightly newer than OGG but with less wide adoption.
- Proprietary Compression – Any non-Werner SF3 SoundFont® compression system. Usage is not allowed in SFe.
- RF64 – See "RIFF64".
- RIFF64 – A 64-bit RIFF-type format. Contrast to the RIFF format, which is 32-bit. Therefore, the maximum file size is above 4 gigabytes in size.
- RIFF-type format – Formats similar to RIFF (Resource Interchange File Format), see "RIFF" in SFSPEC24.PDF for more information.
- ROM samples – Obsolete feature used in legacy sound cards, most modern SF2 files do not use this feature.
- SB – Abbreviation of "Sound Blaster®". For example, "SB X-Fi".
- SFe – A family of enhancements to the SoundFont® 2.04 formats, unofficially created after E-mu/Creative abandoned the original format. May not be structurally compatible with legacy SF2.04.
- SFe 4 – This new specification, based on SoundFont® 2.04 and Werner SF3, with a set of new features making it more realistic. Not to be confused with the incompatible Cognitone SF4 file format.
- SFe-compatible – Indicates files, data, synthesisers, hardware or software that conform to the SFe specification.
- SFe Compression – The compression system based on Werner SF3 that SFe programs should be compliant with.
- Sound Blaster® Live! – The successor to the AWE64, which improved the synthesizer chip to the EMU10K1, supporting modulators.
- Sound Blaster® Audigy – The successor to the SB Live!, containing the EMU10K2 chip.
- Sound Blaster® X-Fi – The successor to the SB Audigy, containing the EMU20K1 or EMU20K2 chip. Supports 24-Bit SoundFont® 2 files (2.04).
- Static RIFF – Any RIFF-type format with a fixed chunk size field width, including RIFF or RIFF64. See "RIFF-type format", "RIFF" and "RIFF64".
- Synth – Abbreviation of "Synthesiser," see "Synthesiser" in SFSPEC24.PDF for more information.
- Tree structure – A structure consisting of branches and leaves.
- Version 3 – See "Werner SF3".
- Version 4 – See "SFe 4".

- Vorbis – A lossy audio compression format commonly used in open-source software. The basic compression format that most Werner SF3 and SFe-compatible software should be expected to implement.
- Werner SF3 – A small upgrade to SoundFont® 2.04 created by Werner Schweer to allow an open source compression solution for SoundFont® programs. Standardised as SFe Compression.

Notes



bring music to new heights